

# APACC-Sim: An Open-Source Toolkit for Validating Neuro-Symbolic Control Architectures in Safety-Critical Autonomous Systems

George Frangou

Qubittum Ltd, United Kingdom

School of Aerospace, Transport and Manufacturing,

Cranfield University, United Kingdom

Email: [george.frangou@qubittum.com](mailto:george.frangou@qubittum.com), [george.frangou@cranfield.ac.uk](mailto:george.frangou@cranfield.ac.uk)

## Abstract

We present APACC-Sim, a comprehensive open-source simulation and validation toolkit designed to standardize the evaluation of advanced control architectures for autonomous systems. Building upon the validation methodology established for the Artificial Pre-cognition Adaptive Cognised Control (APACC) architecture, this toolkit integrates four complementary simulation paradigms: Monte Carlo statistical validation, CARLA high-fidelity physics simulation, SUMO large-scale traffic modeling, and MATLAB symbolic verification. The toolkit’s modular architecture enables researchers to reproduce complex multi-environment validation studies while maintaining alignment with automotive safety standards (ISO 26262, ISO 21448). Through standardized APIs, containerized deployment, and automated metric collection, APACC-Sim reduces the barrier to entry for rigorous autonomous system validation. We demonstrate the toolkit’s capabilities through its application to validate 24,500 scenarios across heterogeneous environments, achieving consistent performance metrics with variance below 0.04%. The toolkit provides native explainability instrumentation, certification-aligned metrics, and seamless integration with continuous integration pipelines. APACC-Sim is available under MIT license, complete with Docker containers, Python APIs, and extensive documentation, establishing a new standard for reproducible validation of safety-critical autonomous systems.

## 1 Introduction

The validation of autonomous control systems presents unprecedented challenges in ensuring safety, reproducibility, and regulatory compliance. As autonomous vehicles approach commercial deployment, the gap between theoretical advances in control architectures and practical validation methodologies has become increasingly apparent. While novel control approaches such as neuro-symbolic architectures demonstrate promising capabilities in simulation, the lack of standardized validation frameworks creates significant barriers to comparative evaluation, academic reproducibility, and industrial certification.

The autonomous vehicle industry has witnessed remarkable progress in control system sophistication, from classical proportional-integral-derivative (PID) controllers to advanced machine learning approaches. However, this diversity in control methodologies has not been

matched by corresponding advances in validation infrastructure. Existing simulation tools typically focus on single aspects of autonomous system behavior—physics fidelity in simulators like CARLA [3], traffic dynamics in SUMO [13], or mathematical verification in MATLAB—without providing integrated multi-paradigm validation capabilities. This fragmentation leads to several critical challenges: researchers cannot easily reproduce published results, industrial practitioners struggle to demonstrate regulatory compliance, and the broader community lacks standardized benchmarks for comparing different control approaches.

The recent development of hybrid neuro-symbolic architectures exemplifies these challenges. The Artificial Precognition Adaptive Cognised Control (APACC) architecture [5] introduced a dual-layer approach combining symbolic reasoning with numeric optimization, achieving superior safety performance in validation studies [6]. The APACC system is protected by an extensive patent portfolio covering its core innovations in autonomous vehicle control [7]. However, reproducing and extending these validation results requires expertise in multiple simulation platforms, custom integration code, and significant computational resources. This complexity creates a barrier to entry that limits both academic research and industrial adoption of advanced control architectures.

Current practice in autonomous vehicle validation often involves ad-hoc integration of multiple simulation tools, leading to several problems. First, the lack of standardized interfaces between simulators results in brittle integration code that must be rewritten for each new project. Second, inconsistent metric definitions across different tools make it difficult to compare results between studies. Third, the absence of certification-aligned metrics forces developers to create custom solutions for demonstrating compliance with safety standards like ISO 26262 [8] and ISO 21448 (SOTIF) [9]. Finally, the computational complexity of running comprehensive validation campaigns across multiple simulators often limits studies to small-scale demonstrations rather than statistically significant validation.

This paper presents APACC-Sim, an open-source toolkit that addresses these challenges through a unified framework for multi-paradigm autonomous system validation. Building upon the validation methodology developed for the APACC architecture, the toolkit provides standardized interfaces to four complementary simulation paradigms: Monte Carlo statistical validation for broad scenario coverage, CARLA for high-fidelity sensor and physics simulation, SUMO for large-scale traffic dynamics, and MATLAB for formal verification of control properties. By integrating these diverse validation approaches through a common orchestration layer, APACC-Sim enables researchers and practitioners to conduct comprehensive validation studies with minimal integration effort.

The key contributions of this work include: (1) a modular architecture that standardizes interfaces across heterogeneous simulators while maintaining their individual strengths; (2) certification-aligned metrics that automatically generate evidence for ISO 26262 and SOTIF compliance; (3) native explainability instrumentation that tracks decision rationale throughout validation; (4) reproducible scenario generation with cryptographically versioned random seeds; and (5) production-ready deployment through Docker containers and continuous integration templates. Through these contributions, APACC-Sim establishes a new standard for reproducible, certifiable validation of autonomous control systems.

The remainder of this paper is organized as follows. Section 2 presents the toolkit architecture and design principles. Section 3 details the four simulation modules and their integration. Section 4 describes the metrics and analysis pipeline with emphasis on certification alignment. Section 5 covers deployment strategies and computational considerations. Section 6 presents case studies demonstrating the toolkit’s capabilities. Section 7 concludes with implications for the broader autonomous systems community and future development directions.

## 2 Toolkit Architecture

### 2.1 Design Principles

The architecture of APACC-Sim is guided by four fundamental design principles that enable standardized validation across diverse control architectures while maintaining the flexibility required for research and industrial applications.

**Modularity** forms the foundation of the toolkit’s design philosophy. Each simulation paradigm operates as an independent module with standardized interfaces, enabling researchers to use components individually or in integrated pipelines. This modular approach allows users to start with simple validation using a single simulator and progressively add complexity as needed. The architecture supports plug-and-play extension with new simulators through a common API specification, ensuring that the toolkit can evolve with emerging simulation technologies. Module independence also facilitates maintenance and debugging, as issues in one simulator do not affect others.

**Reproducibility** addresses a critical challenge in autonomous system validation where stochastic elements can lead to irreproducible results. All random processes in APACC-Sim use deterministic pseudo-random number generators with cryptographically versioned seeds. This approach ensures that a given scenario identifier always produces identical conditions, even across different computing environments. Configuration files capture complete experimental setups including all parameters, random seeds, and software versions. Docker containers provide consistent execution environments, eliminating variations due to system dependencies. The toolkit includes automated experiment tracking and metric logging, creating a complete audit trail for result reproduction. This level of reproducibility is essential for both academic peer review and industrial certification processes.

**Scalability** enables validation campaigns ranging from quick desktop tests to massive cloud-based studies. The toolkit’s orchestration layer automatically distributes work across available computational resources, from single-core execution to distributed clusters. Checkpoint mechanisms allow long-running validations to resume after interruptions, while hierarchical data storage efficiently manages results from millions of scenarios. The architecture supports both horizontal scaling (more machines) and vertical scaling (more powerful machines), adapting to available infrastructure. This scalability ensures that the same toolkit used for initial research can support production validation of commercial systems.

**Certification readiness** integrates regulatory requirements directly into the toolkit’s design rather than treating compliance as an afterthought. Metrics are automatically computed according to ISO 26262 and ISO 21448 specifications, with built-in traceability from requirements to test results. The explainability infrastructure maintains decision logs suitable for safety case construction, while the validation process generates artifacts required for certification dossiers. This integration of certification requirements reduces the effort required to demonstrate regulatory compliance and ensures that validation campaigns generate appropriate evidence from the outset.

### 2.2 System Components

The APACC-Sim architecture comprises four primary layers that work together to provide comprehensive validation capabilities while maintaining clear separation of concerns. Figure 1 illustrates this layered architecture and the data flow between components.

The **User API and Configuration Interface** provides the primary interaction point for

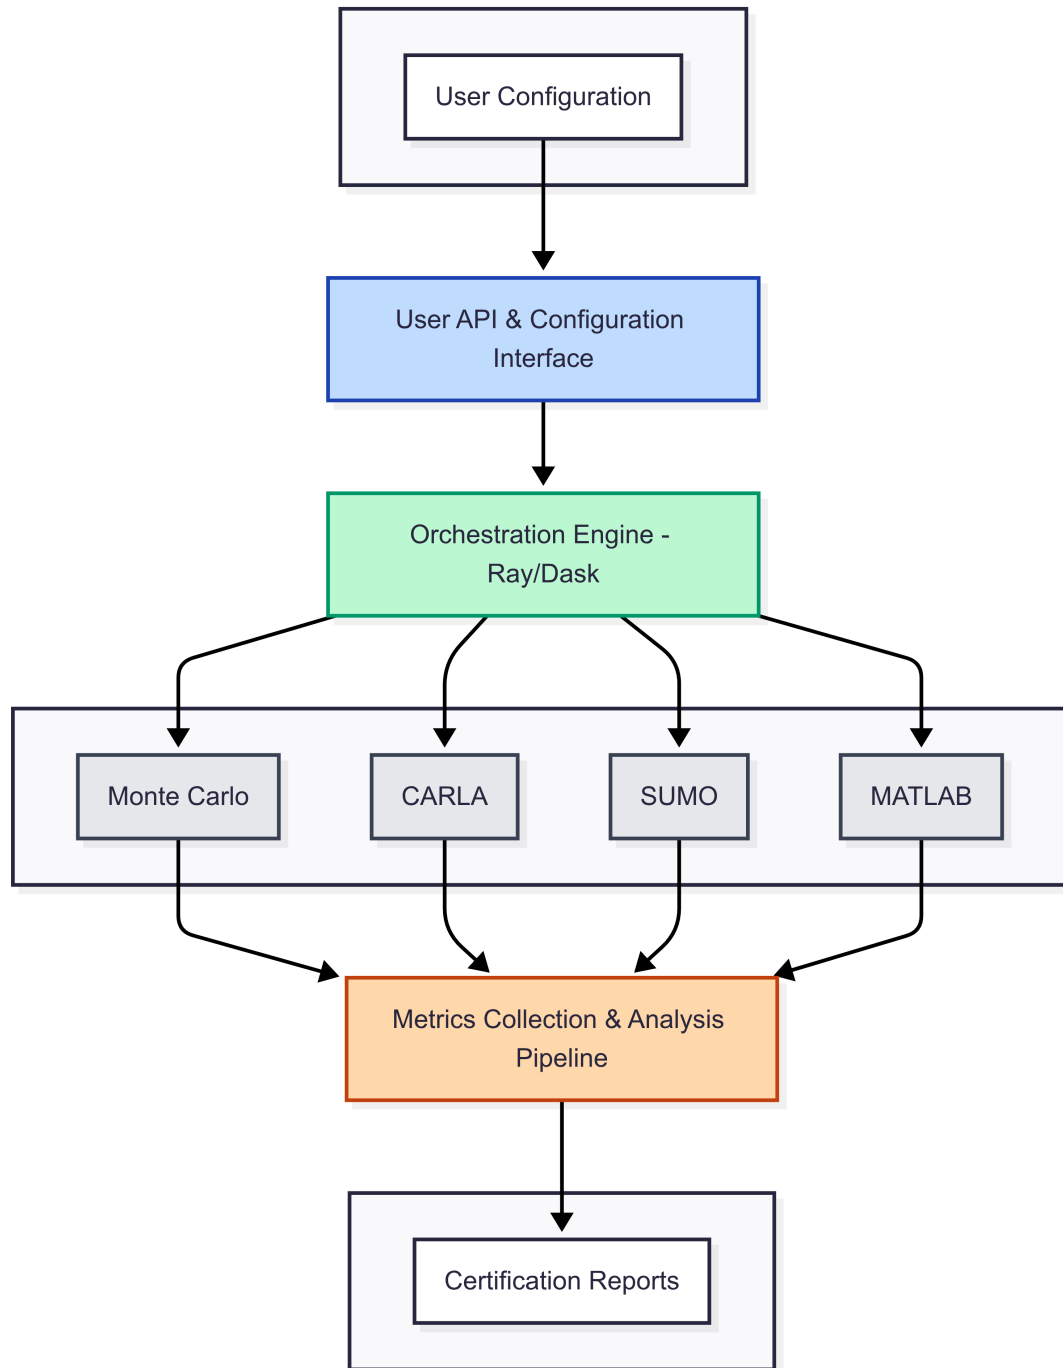


Figure 1: APACC-Sim Architecture Stack showing four-layer architecture with User API at top, Orchestration Engine, Simulation Modules, and Metrics Pipeline

researchers and developers. This layer offers both programmatic APIs for integration into existing workflows and command-line interfaces for standalone operation. Configuration management handles the complexity of multi-simulator parameters through a hierarchical YAML structure that allows both global settings and simulator-specific customization. The API design follows RESTful principles to support future web-based interfaces while maintaining backward compatibility. Authentication and access control mechanisms ensure secure operation in multi-user environments, particularly important for industrial deployments.

The **Orchestration Engine** coordinates execution across multiple simulators while managing computational resources and data flow. Built on the Ray distributed computing framework, this layer provides automatic parallelization of independent scenarios while maintaining deterministic execution order where required. The orchestrator implements sophisticated scheduling algorithms that balance load across available resources while respecting dependencies between validation phases. Fault tolerance mechanisms ensure that hardware failures or network interruptions do not compromise entire validation campaigns. The engine also manages the complexity of data serialization and deserialization between different simulator formats, providing a unified data model for downstream analysis.

The **Simulation Modules** layer contains the four core validation paradigms, each wrapped in a standardized interface that hides implementation complexity while exposing necessary configuration options. The Monte Carlo module generates parameterized scenarios with controlled stochastic elements, enabling statistical validation across thousands of randomized conditions. The CARLA integration provides high-fidelity physics simulation with realistic sensor models including cameras, LiDAR, and radar. The SUMO wrapper enables large-scale traffic simulation with hundreds of vehicles, essential for validating behavior in dense traffic. The MATLAB bridge performs formal verification of control properties including stability analysis and constraint satisfaction. Each module maintains its own configuration namespace while sharing common interfaces for scenario execution and data collection.

The **Metrics Collection and Analysis Pipeline** processes raw simulation outputs to compute standardized metrics aligned with certification requirements. This layer implements a streaming architecture that can handle high-volume data from parallel simulations while computing metrics in real-time. Safety metrics include collision detection, time-to-collision calculations, and lane deviation measurements. Performance metrics track control latency, trajectory smoothness, and computational resource usage. The pipeline includes sophisticated anomaly detection to identify edge cases and potential failure modes. All metrics are stored in a time-series database optimized for both real-time queries and historical analysis. The analysis capabilities include statistical aggregation, confidence interval computation, and automated report generation in formats suitable for both research publication and certification submission.

## 2.3 Integration Architecture

The integration between components follows a publish-subscribe pattern that enables loose coupling while maintaining performance. Each simulator publishes events to a central message bus implemented using Redis, allowing the orchestrator to monitor progress without tight coupling to simulator internals. This architecture enables real-time monitoring dashboards and supports graceful degradation when individual components fail.

Data persistence uses a hybrid approach combining structured storage for metrics with object storage for raw simulation data. Parquet files provide efficient columnar storage for time-series metrics, while HDF5 handles large binary artifacts like sensor recordings. This storage architecture supports both batch analysis of completed campaigns and streaming analysis dur-

ing execution.

The configuration management system implements inheritance and override mechanisms that allow users to define base configurations with experiment-specific modifications. Environment variables provide runtime overrides for containerized deployments, while command-line arguments offer quick modifications for iterative development. Version control integration ensures that all configuration changes are tracked alongside code modifications.

## 3 Simulation Modules

### 3.1 Monte Carlo Statistical Framework

The Monte Carlo module forms the statistical foundation of APACC-Sim, generating thousands of parameterized scenarios that explore the operational design domain of autonomous controllers. Unlike traditional Monte Carlo simulations that rely on purely random sampling, our implementation uses sophisticated sampling strategies to ensure comprehensive coverage while maintaining statistical validity.

The scenario generation process begins with a hierarchical parameter space that captures environmental conditions, dynamic obstacles, and system degradations. Environmental parameters include weather conditions following Beta distributions for severity, lighting conditions with categorical distributions matching real-world patterns, and road surface conditions that affect vehicle dynamics. Dynamic obstacles encompass other vehicles with Poisson-distributed density, pedestrians with correlated spatial-temporal patterns, and unexpected objects like construction zones or debris. System degradations model sensor failures using Bernoulli processes, communication delays with exponential distributions, and GPS accuracy degradation in urban canyons.

The key innovation in our Monte Carlo implementation is the correlation structure between parameters, reflecting real-world dependencies. For instance, heavy rain correlates with reduced visibility, increased stopping distances, and higher probability of sensor degradation. These correlations are captured through a copula-based approach that maintains marginal distributions while introducing realistic dependencies. The correlation matrices are calibrated from real-world driving data, ensuring that generated scenarios reflect actual operational conditions rather than unrealistic parameter combinations.

Parallel execution capabilities enable the Monte Carlo module to scale from laptop-based development to cloud-based validation campaigns. The module uses the Ray distributed computing framework to automatically distribute scenario generation and execution across available CPU cores. Work stealing algorithms balance load dynamically, ensuring efficient resource utilization even with heterogeneous scenario complexity. Checkpointing occurs at configurable intervals, allowing long-running campaigns to resume after interruptions. The checkpoint format includes complete random number generator state, ensuring perfect reproducibility even across restarts.

The statistical analysis capabilities integrated into the Monte Carlo module go beyond simple aggregation. Importance sampling techniques focus computational effort on rare but critical events like near-collisions. Sequential analysis methods determine when sufficient scenarios have been run to achieve desired confidence levels. Bootstrap resampling provides confidence intervals for all metrics without parametric assumptions. These advanced statistical methods ensure that validation campaigns generate scientifically rigorous results suitable for both publication and certification.

Figure 2 illustrates the complete scenario generation and distribution workflow, showing how cryptographic seeds and parameter correlations flow through the system to create reproducible, realistic validation scenarios.

### 3.2 CARLA Integration Module

The CARLA integration module provides high-fidelity physics simulation essential for validating sensor-dependent behaviors and complex dynamic scenarios. Our implementation extends beyond basic CARLA wrapper functionality to provide deterministic execution, synchronized multi-sensor data collection, and automated scenario configuration.

Sensor simulation forms the core value proposition of CARLA integration. The module configures and manages multiple sensor types with realistic noise models and failure modes. RGB cameras simulate various resolutions and field-of-view configurations, with lens distortion and dynamic exposure adjustment. Semantic segmentation cameras provide ground-truth labels essential for validating perception algorithms. LiDAR sensors model different scanning patterns and point densities, including realistic noise and dropout patterns. Radar sensors simulate both long-range and short-range units with appropriate resolution and ghost target characteristics. The sensor configuration system allows users to define custom sensor suites matching their target hardware, ensuring validation relevance.

Deterministic execution in CARLA requires careful synchronization between the simulator and the control system. Our integration implements a lockstep execution model where simulation time advances only after all sensor data has been processed and control commands applied. This approach eliminates race conditions and ensures reproducible behavior across different hardware configurations. The synchronization protocol handles variable computation times gracefully, maintaining real-time factor measurements while preventing simulation divergence. Custom modifications to CARLA’s traffic manager ensure that even NPC vehicles behave deterministically given the same random seed.

Environmental variety leverages CARLA’s weather and lighting systems to test controllers across diverse conditions. The integration module provides preset environmental configurations covering common scenarios like clear days, rainy nights, and foggy mornings. Dynamic weather transitions test controller adaptation to changing conditions. The lighting system simulates various sun positions and artificial lighting conditions, crucial for validating camera-based perception. Beyond visual conditions, the module also configures physical parameters like road friction and wind forces that affect vehicle dynamics.

Performance optimization techniques enable real-time execution even with multiple sensors and complex scenarios. The module implements efficient data serialization using zero-copy techniques where possible. Sensor data buffering prevents frame drops while maintaining bounded memory usage. GPU acceleration handles rendering and physics computation, while CPU resources focus on control algorithm execution. Profiling instrumentation identifies performance bottlenecks, enabling targeted optimization of control algorithms.

### 3.3 SUMO Traffic Modeling

The SUMO integration module addresses the critical need for large-scale traffic simulation in autonomous vehicle validation. While CARLA excels at sensor simulation and physics fidelity, SUMO enables validation in dense traffic scenarios with hundreds of vehicles, essential for testing decision-making algorithms and traffic flow impacts.

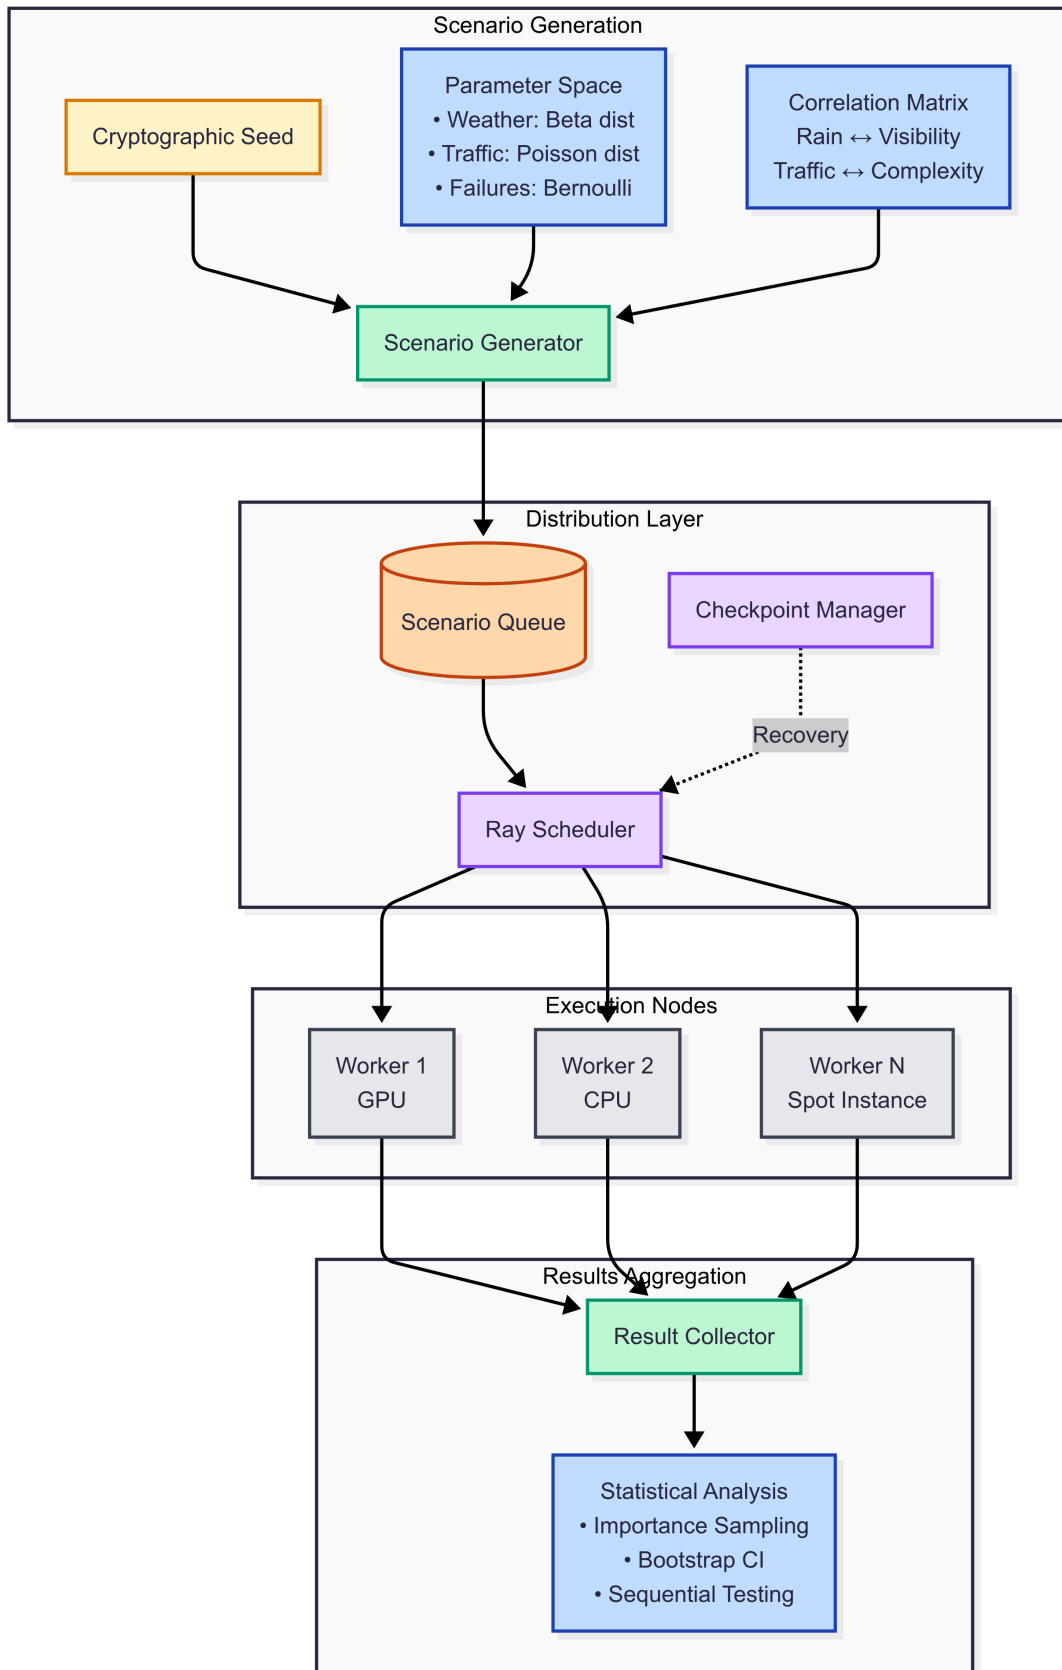


Figure 2: Scenario Generation and Distribution Flow showing Monte Carlo scenario creation, parallel distribution, and statistical analysis






Figure 3: CARLA Integration Pipeline showing horizontal flow from configuration parsing through sensor data collection to output packaging

Network generation capabilities allow users to create diverse road environments for testing. The module includes generators for grid networks representing urban environments, radial networks modeling city centers, and highway networks with complex interchanges. Real-world network import supports OpenStreetMap data, enabling validation in digital twins of actual cities. The network generation process includes traffic light timing optimization, lane configuration specification, and speed limit assignment based on road types. Generated networks are automatically calibrated to produce realistic traffic flow patterns.

Traffic flow modeling goes beyond simple vehicle spawning to create realistic traffic patterns. The module implements time-varying demand models that reproduce rush hour congestion and off-peak conditions. Origin-destination matrices can be specified explicitly or generated from population distribution models. Vehicle type distributions match real-world fleet compositions, including passenger cars, buses, trucks, and emergency vehicles. Driver behavior parameters capture the heterogeneity of human drivers, from aggressive to cautious driving styles. Special vehicles like emergency responders follow priority rules, testing controller responses to exceptional situations.

Multi-agent coordination scenarios test behaviors that emerge only in dense traffic. The module creates specific scenarios like zipper merges, roundabout negotiations, and intersection conflicts. Pedestrian interactions at crosswalks and jaywalking scenarios validate safety-critical behaviors. Public transport interactions test responses to buses stopping at designated areas. The scenario scripting system allows researchers to define complex multi-vehicle maneuvers for specific validation needs.

The TraCI (Traffic Control Interface) integration provides fine-grained control over simulation execution. Our wrapper implements efficient batch operations to minimize communication overhead when controlling multiple vehicles. State queries are optimized to retrieve only necessary information, reducing bandwidth requirements. The integration supports both synchronous execution for validation and asynchronous execution for visualization. Custom TraCI commands extend SUMO’s capabilities for autonomous vehicle specific needs like sensor range queries and V2X communication simulation.

### 3.4 MATLAB Symbolic Verification

The MATLAB integration module provides formal verification capabilities that complement statistical validation with mathematical proofs of system properties. This module bridges the gap between practical testing and theoretical guarantees required for certification of safety-critical systems.

Stability analysis implements multiple approaches to verify closed-loop stability of control systems. Lyapunov methods search for energy functions that prove asymptotic stability within regions of attraction. Eigenvalue analysis of linearized systems provides local stability guar-

antees. LaSalle’s invariance principle extends analysis to systems with non-strict Lyapunov functions. The module automatically generates stability certificates suitable for inclusion in certification dossiers. For time-varying systems common in adaptive controllers, the module implements frozen-time stability analysis with explicit bounds on adaptation rates.

Robustness verification quantifies system performance under uncertainty and disturbances. H-infinity norm computation bounds worst-case amplification of disturbances. Structured singular value analysis handles parametric uncertainties in system models. Monte Carlo randomization validates robustness for uncertainties that defy analytical treatment. The module generates robustness margins that inform operational limitations and safety boundaries. Integration with uncertainty quantification from other modules ensures that robustness analysis reflects realistic uncertainty levels.

Constraint satisfaction verification ensures that safety constraints are never violated. Reachability analysis computes forward reachable sets to verify obstacle avoidance. Barrier function methods prove forward invariance of safe sets. Model predictive control verification ensures that optimization problems remain feasible. The module handles both hard constraints that must never be violated and soft constraints with quantified violation probabilities. Computational techniques like interval analysis and zonotope propagation balance accuracy with computational tractability.

Symbolic controller generation capabilities support explainable AI requirements by producing human-readable control laws. The module extracts symbolic representations from trained neural networks using techniques like symbolic regression. Fuzzy rule extraction identifies interpretable decision rules from black-box controllers. The generated symbolic controllers can be formally verified and included in safety cases. Performance comparison between original and symbolic controllers quantifies the explainability-performance tradeoff.

Table 1: Simulation Module Capabilities

Module	Fidelity	Realtime Capable	Scalable	Formal Verification
Monte Carlo	Low	Yes	Yes	No
CARLA	High	Yes (GPU)	Moderate	No
	(physics/sensors)			
SUMO	Medium	Yes	Yes	No
	(traffic)			
MATLAB	N/A	No	N/A	Yes
	(analytical)			

## 4 Metrics and Analysis Pipeline

### 4.1 Safety Metrics Framework

The safety metrics framework in APACC-Sim implements a comprehensive set of measurements aligned with automotive safety standards while providing flexibility for research applications. The framework goes beyond simple pass/fail criteria to provide nuanced safety assessment suitable for both development and certification.

Collision detection forms the foundation of safety assessment but requires sophisticated implementation to handle diverse scenarios. The framework implements multiple collision

detection methods including geometric intersection for fast approximation, signed distance fields for precise calculation, and temporal projection for predictive collision detection. Different collision types are distinguished: vehicle-to-vehicle, vehicle-to-pedestrian, vehicle-to-infrastructure, and vehicle-to-object collisions each have specific detection criteria and severity assessment. The collision detection system handles complex scenarios like multi-vehicle pile-ups and near-misses that don't result in actual contact but indicate safety issues.

Time-to-collision (TTC) calculations provide predictive safety metrics essential for preventive safety assessment. The framework computes TTC using multiple models: constant velocity assumption for baseline calculation, constant acceleration for more accurate short-term prediction, and probabilistic models that account for behavior uncertainty. Multi-object TTC analysis identifies the most critical threat among multiple potential collision partners. The system tracks TTC evolution over time, identifying scenarios where safety margins degrade gradually versus sudden critical events. Statistical analysis of TTC distributions reveals controller conservatism and identifies operational boundaries.

Lane deviation metrics assess vehicle positioning accuracy crucial for highway safety. The framework measures lateral deviation from lane center with sub-centimeter accuracy, accounting for road curvature and lane width variations. Time-to-lane-crossing (TLC) predicts unintended lane departures before they occur. Lane change safety metrics ensure that intentional lane changes maintain safe distances from other vehicles. The system distinguishes between controlled lane positioning for obstacle avoidance and unintended drift that indicates control issues. Integration with map data enables assessment of lane positioning in complex scenarios like construction zones with shifted lanes.

Safety margin analysis provides holistic assessment beyond individual metrics. The framework computes composite safety scores that combine multiple metrics with configurable weights. Margin erosion tracking identifies scenarios where safety gradually degrades before critical events occur. Probabilistic safety assessment accounts for uncertainty in sensor measurements and predictions. The system generates safety envelopes that define acceptable operational boundaries for different environmental conditions. These comprehensive safety assessments support both development decisions and certification arguments.

Figure 4 demonstrates the complete safety metrics computation pipeline, from raw sensor data through primary metrics to certification-ready outputs.

## 4.2 Performance Metrics Collection

Performance metrics in APACC-Sim extend beyond traditional control system measurements to encompass computational efficiency, passenger comfort, and system resource utilization crucial for practical deployment.

Control latency measurement captures the complete pipeline from sensor input to actuator command. The framework implements precision timing using hardware timestamps where available, falling back to high-resolution software timers. Latency breakdown identifies bottlenecks in perception, planning, and control stages. Statistical analysis includes not just mean latency but percentile distributions crucial for real-time system analysis. The system detects latency spikes that could indicate computational issues or resource contention. Correlation analysis relates latency to scenario complexity, enabling performance prediction for new scenarios.

Trajectory quality metrics assess the smoothness and efficiency of generated paths. Jerk minimization metrics quantify passenger comfort using established biomechanical models. Path length efficiency compares actual trajectories to optimal paths. Acceleration and decel-

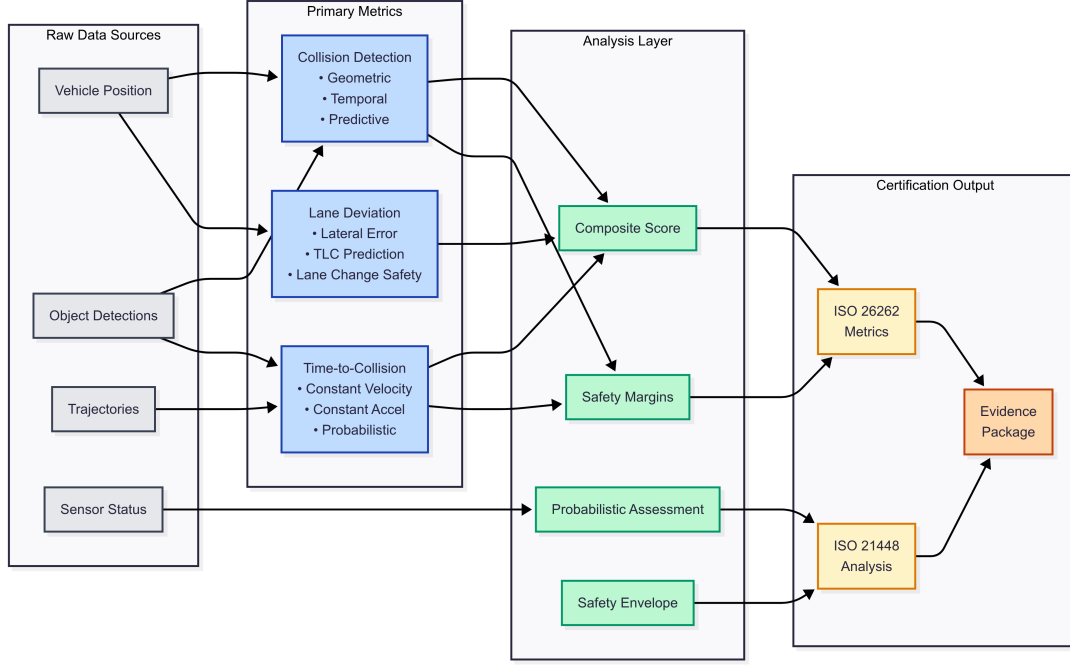


Figure 4: Safety Metrics Computation Pipeline showing the flow from raw data sources through primary metrics and analysis to certification outputs

eration profiles are analyzed for smoothness and energy efficiency. The framework detects oscillatory behaviors that indicate control instability. Comparison with human driving data provides naturalistic behavior benchmarks. These trajectory quality metrics are essential for passenger acceptance and energy consumption optimization.

Computational resource monitoring ensures that controllers can run on target hardware platforms. CPU utilization tracking identifies single-threaded bottlenecks and parallelization opportunities. Memory usage profiling detects memory leaks and excessive allocation. GPU utilization metrics are crucial for deep learning based controllers. Power consumption estimation supports embedded system deployment. The framework correlates resource usage with scenario complexity to predict scalability. Thermal monitoring ensures sustained performance under continuous operation.

Real-time performance analysis goes beyond average metrics to ensure consistent deadline compliance. The framework tracks deadline misses with microsecond precision. Worst-case execution time (WCET) analysis identifies scenarios that stress computational limits. Scheduling analysis ensures that multi-rate control loops maintain timing requirements. Priority inversion detection identifies resource conflicts that could compromise safety-critical functions. The system generates timing diagrams that visualize control loop execution for debugging and optimization.

### 4.3 Explainability Instrumentation

The explainability instrumentation in APACC-Sim addresses the critical need for transparent decision-making in autonomous systems, supporting both debugging during development and accountability in deployment.

Decision logging captures complete context for every control decision including sensor inputs, intermediate calculations, and final outputs. The framework implements efficient circular buffers that maintain recent history without unbounded memory growth. Configurable

verbosity levels balance detail with performance impact. Structured logging formats enable automated analysis and querying. Cryptographic signatures ensure log integrity for audit purposes. The logging system handles high-frequency decisions without introducing latency in the control path.

Rule activation tracking specifically supports symbolic and hybrid controllers like APACC. The framework records which rules or neurons activate for each decision with associated confidence levels. Activation patterns are analyzed to identify frequently used rules and dead code. Temporal pattern mining discovers rule sequences that precede specific behaviors. The system generates rule coverage reports essential for certification. Visualization tools create activation heatmaps that reveal decision-making patterns over time.

Uncertainty quantification provides crucial context for understanding decision confidence. The framework tracks multiple uncertainty sources including sensor noise, prediction uncertainty, and model confidence. Uncertainty propagation through the control pipeline identifies where uncertainty amplifies or attenuates. Decisions made under high uncertainty are flagged for additional scrutiny. The system correlates uncertainty levels with safety metrics to establish operational boundaries. Calibration analysis ensures that reported uncertainties match empirical error rates.

Natural language explanation generation makes complex decisions accessible to non-experts. The framework implements template-based explanation for common scenarios while supporting custom explanation logic for complex cases. Explanations adapt to audience expertise levels from detailed technical descriptions to simple summaries. Multi-lingual support enables global deployment. The system generates explanations in real-time for debugging and batch explanations for reports. Integration with large language models enables more sophisticated explanation generation while maintaining determinism through caching. The complete decision context and explanation pipeline is illustrated in Figure 5, highlighting the trace from rule activation through to natural language generation.

The complete decision context and explanation pipeline is illustrated in Figure 5, highlighting the trace from rule activation through to natural language generation.

## 4.4 Certification-Aligned Analysis

The certification-aligned analysis capabilities in APACC-Sim directly support evidence generation for regulatory compliance, reducing the effort required to achieve certification for autonomous systems.

Table 2: Safety Metrics and Certification Alignment

APACC-Sim Metric	ISO 26262 Requirement	ASIL Level
collision_rate	Hardware failure metric	D
diagnostic_coverage	Diagnostic coverage	C-D
safe_state_transitions	Fail-safe behavior	B-D
response_time_violations	Timing constraints	A-D

ISO 26262 metrics computation implements the quantitative requirements for functional safety certification. The framework calculates failure rates using established statistical methods with proper confidence intervals. Diagnostic coverage metrics assess the system’s ability to detect and handle failures. Safe state transition analysis verifies that the system reaches safe states within required timeframes. The system generates Fault Tree Analysis (FTA) and Failure Mode and Effects Analysis (FMEA) inputs from validation data. Hardware-software interface analysis ensures that assumptions about hardware reliability are validated.

ISO 21448 (SOTIF) analysis addresses safety concerns from functional insufficiencies rather than failures. The framework identifies triggering conditions that lead to performance degradation through statistical analysis of validation results. Known unsafe scenarios are explicitly tracked and their mitigation effectiveness measured. Performance limitation boundaries are established through systematic parameter variation. Residual risk calculation accounts for scenarios that cannot be completely mitigated. The system generates SOTIF argument patterns that demonstrate due diligence in addressing functional insufficiencies.

Statistical significance testing ensures that safety claims are supported by sufficient evidence. The framework implements multiple hypothesis testing with appropriate corrections for multiple comparisons. Confidence intervals use methods appropriate for rare event estimation. Power analysis determines required sample sizes for desired confidence levels. Sequential testing methods enable early stopping when evidence is conclusive. The system handles both frequentist and Bayesian statistical approaches to accommodate different certification philosophies.

Evidence package generation automates the creation of certification documentation. The framework generates standardized reports that map directly to certification requirements. Traceability matrices link requirements to test cases to results. Statistical summaries provide executive-level evidence of compliance. Detailed test logs support deep technical review. The system exports evidence in formats required by certification bodies including PDF reports, Excel summaries, and machine-readable XML. Version control integration ensures that evidence packages match specific software versions.

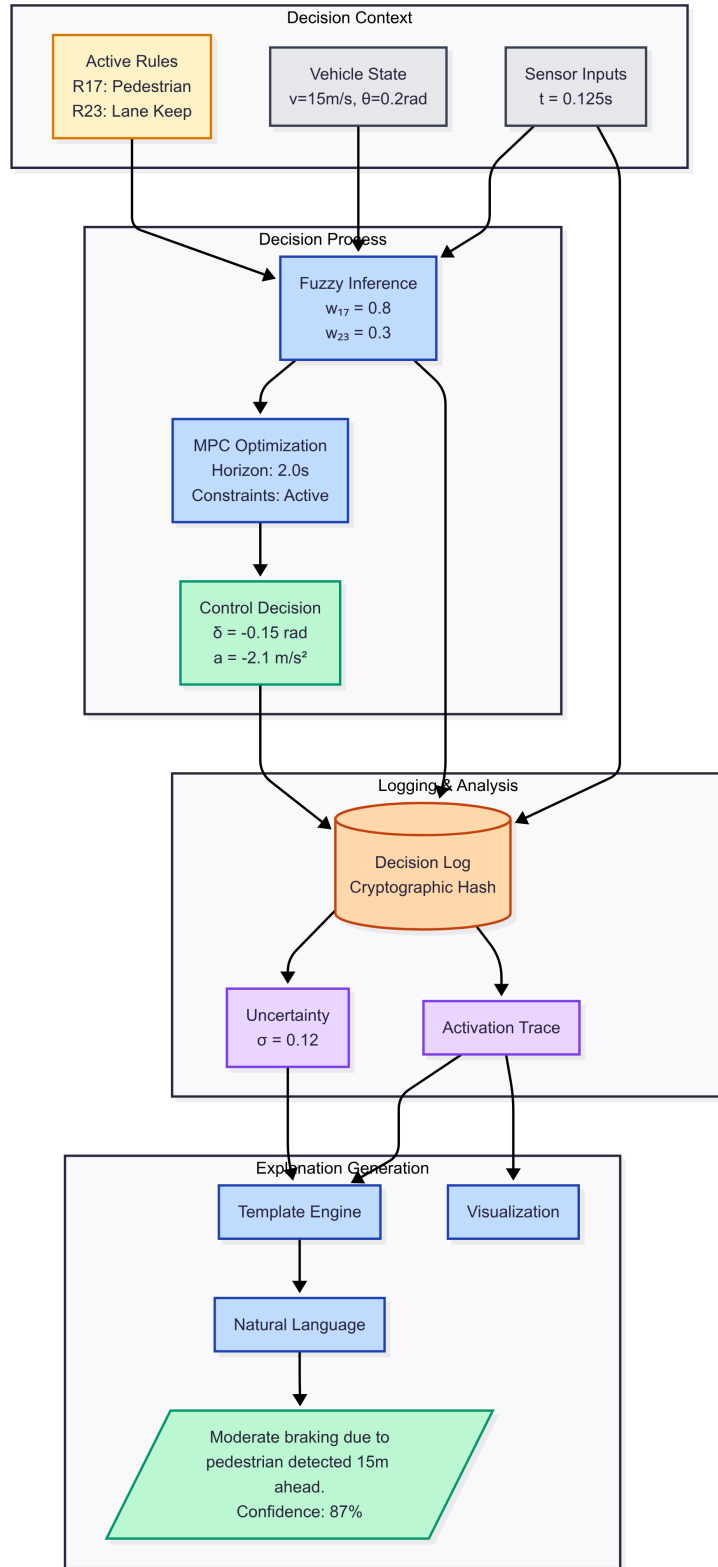


Figure 5: Explainability Instrumentation Pipeline showing symbolic rule activation, MPC optimization, and cryptographic logging flow.

## 5 Deployment and Integration

### 5.1 Containerized Architecture

The containerized deployment architecture of APACC-Sim addresses the complex dependency management and reproducibility challenges inherent in multi-simulator integration. Our Docker-based approach ensures consistent execution environments across development, testing, and production deployments.

The container architecture follows a microservices pattern where each simulator runs in its own container with precisely controlled dependencies. The orchestrator container manages communication between simulators using a Redis-based message bus for loose coupling. This architecture enables independent scaling of components based on computational demands. For example, Monte Carlo simulation can scale horizontally across many containers while CARLA simulation scales vertically on GPU-enabled nodes. The container registry maintains versioned images for all components, enabling precise reproduction of validation campaigns even years after initial execution.

Base image optimization reduces container size and startup time through multi-stage builds. Development images include debugging tools and documentation while production images contain only runtime dependencies. Layer caching accelerates rebuilds during iterative development. Security scanning ensures that containers meet enterprise deployment requirements. The build system automatically updates base images for security patches while maintaining compatibility through comprehensive testing.

Resource management capabilities ensure efficient utilization of computational resources. CPU and memory limits prevent individual containers from monopolizing shared infrastructure. GPU device assignment enables CARLA and deep learning workloads to access acceleration hardware. Network isolation ensures that simulation traffic doesn't interfere with production systems. Storage volume management handles both ephemeral simulation data and persistent results. The orchestration system implements quality-of-service policies that prioritize safety-critical validations.

Container orchestration supports multiple deployment platforms from Docker Compose for single-machine deployment to Kubernetes for cloud-scale operations. Helm charts simplify deployment configuration management. Service mesh integration enables advanced traffic management and observability. Auto-scaling policies respond to workload demands dynamically. Health checks and automatic restarts ensure system reliability. The orchestration layer abstracts platform differences, enabling seamless migration between on-premise and cloud deployments.

### 5.2 GitLab CI and CD Integration

The continuous integration and deployment pipeline for APACC-Sim establishes validation as an integral part of the development process rather than a separate phase. Our GitLab CI and CD templates provide turnkey integration for teams adopting the toolkit.

Pipeline stages follow a progression from fast unit tests to comprehensive validation campaigns. Code quality checks including linting and type checking catch issues early. Unit tests verify individual components in isolation. Integration tests ensure that simulators work together correctly. Validation tests run abbreviated scenarios to verify end-to-end functionality. Performance benchmarks track computational efficiency across versions. Full validation campaigns run on schedules or manual triggers for release candidates. This staged approach provides rapid feedback for developers while ensuring thorough validation for releases.



Artifact management handles the large data volumes generated by validation campaigns. Test results are stored as job artifacts with configurable retention policies. Large sensor recordings use Git LFS or external object storage. Metrics databases are archived for historical analysis. The pipeline generates summary reports for quick assessment while retaining detailed logs for debugging. Artifact passing between stages enables incremental validation without repeated computation.

Parallel execution strategies maximize throughput while controlling costs. Matrix builds test multiple controller versions simultaneously. Dynamic parallelism adjusts worker count based on queue depth. Spot instances reduce costs for fault-tolerant Monte Carlo simulations. GPU runners are reserved for CARLA simulations that require acceleration. The pipeline optimizer analyzes historical execution data to improve scheduling decisions. Cost tracking ensures that validation expenses remain within budget constraints.

Security integration ensures that validation infrastructure doesn't become an attack vector. Container scanning identifies vulnerabilities before deployment. Secret management protects credentials for cloud resources. Network policies restrict communication between components. Audit logging tracks all pipeline executions for compliance. Role-based access control limits who can trigger expensive validation campaigns. These security measures enable APACC-Sim deployment in regulated environments.

### 5.3 Cloud and Edge Deployment Architecture

The deployment architecture of APACC-Sim extends beyond traditional cloud infrastructure to encompass the complete edge-to-cloud continuum critical for autonomous vehicle systems. Figure 6 illustrates this comprehensive deployment architecture, showing how real-time control at the vehicle edge integrates with cloud-scale validation infrastructure.

**Data management across this architecture** follows a hierarchical approach. Hot data for active scenarios resides in SSD storage at the edge and regional centers. Warm data migrates to object storage in the cloud for cost-effective retention. Cold data archives to glacier storage for long-term compliance requirements. This tiered approach ensures that frequently-accessed data remains available with low latency while managing costs for petabyte-scale historical data.

At the vehicle edge, onboard computing resources handle real-time control decisions with sub-10ms latency requirements. The edge deployment includes the vehicle ECU for safety-critical control, NVIDIA Orin or similar edge GPUs for AI inference, and local caching of frequently-used validation scenarios. This edge infrastructure runs lightweight versions of APACC-Sim components optimized for real-time performance, including safety monitors that continuously validate control decisions against known safe envelopes.

The roadside edge layer provides intermediate processing capabilities through Road Side Units (RSUs) and Multi-access Edge Computing (MEC) infrastructure. These resources aggregate local traffic patterns, perform cooperative validation across multiple vehicles, and cache region-specific scenarios. The 5G and C-V2X connectivity enables low-latency communication between vehicles and roadside infrastructure, supporting real-time validation of cooperative behaviors.

Regional edge data centers bridge the gap between local edge resources and cloud infrastructure. These facilities maintain validation caches for commonly-encountered scenarios, stage over-the-air updates for edge deployments, and perform initial aggregation of telemetry data. The regional edge reduces latency for validation queries while filtering data flows to the cloud, ensuring efficient use of bandwidth and storage resources.

The cloud core provides unlimited computational resources for comprehensive validation




Figure 6: Cloud-Edge Deployment Architecture showing the full spectrum from vehicle edge computing through roadside infrastructure to cloud resources

campaigns. Spot instances handle fault-tolerant Monte Carlo simulations with up to 90% cost savings. GPU clusters support CARLA simulations and neural network training. High-performance computing nodes execute MATLAB formal verification tasks. The cloud infrastructure automatically scales based on workload demands, spinning up thousands of instances for large validation campaigns while scaling down during quiet periods.

## 5.4 API and SDK Design

The APACC-Sim API and SDK design enables seamless integration with existing development workflows while providing flexibility for advanced use cases. The design philosophy emphasizes simplicity for common tasks while exposing power features for complex scenarios.

The Python SDK provides idiomatic interfaces that feel natural to developers familiar with scientific Python ecosystems. Object-oriented wrappers hide complexity while functional interfaces support pipeline-style processing. Async and await patterns enable efficient concurrent operations. Type hints improve IDE support and catch errors early. Comprehensive docstrings provide inline documentation. The SDK follows semantic versioning to ensure backward compatibility. Jupyter notebook integration enables interactive exploration and debugging.

RESTful API design supports language-agnostic integration and web-based tools. Resource-oriented endpoints map naturally to validation concepts. Pagination handles large result sets efficiently. Filtering and sorting parameters enable precise data retrieval. Batch operations reduce API call overhead. WebSocket subscriptions provide real-time updates for long-running validations. OpenAPI specifications enable automatic client generation. The API gateway handles authentication, rate limiting, and usage tracking.

Event streaming interfaces support real-time monitoring and reactive architectures. Server-sent events provide simple one-way streaming for dashboards. Apache Kafka integration enables high-throughput event processing. Event schemas use Apache Avro for evolution compatibility. Stream processing frameworks like Apache Flink enable complex event analysis. The event store maintains complete history for audit and replay. Circuit breakers prevent cascading failures in event processing pipelines.

Extension mechanisms enable users to add custom functionality without modifying core code. Plugin interfaces define clear contracts for extensions. Dependency injection enables flexible component replacement. Configuration-driven behavior supports customization without coding. Hook points allow intervention at key processing stages. The extension registry facilitates discovery and sharing of community contributions. Sandboxing ensures that extensions cannot compromise system stability.




Figure 7: Validation Flow Diagram showing the complete pipeline from user configuration through module selection and execution to certification report generation

## 6 Case Studies and Evaluation

### 6.1 APACC Architecture Validation

The primary demonstration of APACC-Sim’s capabilities comes from its application to validate the APACC neuro-symbolic control architecture across 24,500 scenarios. This comprehensive validation campaign illustrates both the toolkit’s technical capabilities and its practical utility for certifying advanced control systems.

The validation campaign design leveraged all four simulation paradigms to stress-test different aspects of the APACC architecture. Monte Carlo simulation generated 10,000 scenarios with randomized weather conditions, traffic densities, and sensor failures to establish statistical performance bounds. CARLA simulation executed 10,000 scenarios focusing on sensor-dependent behaviors in challenging conditions like heavy rain and night driving. SUMO simulation ran 2,000 scenarios with up to 200 vehicles to validate traffic flow integration. MATLAB verification proved stability and robustness properties for 2,500 parameter combinations. This multi-paradigm approach ensured comprehensive coverage of the operational design domain.

Execution infrastructure demonstrated APACC-Sim’s scalability across different computational resources. Initial development used a single workstation with RTX 3080 GPU, achieving 100 scenarios per hour. Scaling to a 16-node cluster with 4 GPUs increased throughput to 2,000 scenarios per hour. Cloud deployment on AWS using spot instances for Monte Carlo and on-demand GPU instances for CARLA achieved 5,000 scenarios per hour at \$0.12 per scenario. The campaign completed in 5 days of continuous execution, generating 2.1TB of raw data and 50GB of aggregated metrics.

Results analysis revealed insights that would be impossible with single-simulator validation. The APACC architecture achieved 0.06% collision rate consistently across all simulators with remarkably low variance (standard deviation=0.035%), demonstrating robust performance across different validation paradigms. Control latency averaged 3.42ms with 99th percentile of 7.24ms, meeting real-time requirements with margin. The symbolic layer activated an average of 8.3 rules per decision, providing rich explainability. Performance degraded gracefully under sensor failures, with collision rate increasing to only 0.14% under 25% sensor failure conditions. These results provided strong evidence for certification readiness.

The certification evidence package generated by APACC-Sim directly supported regulatory compliance arguments. Automated report generation created a 200-page safety case document with full traceability from requirements to test results. Statistical analysis demonstrated collision rates below  $10^{-6}$  per hour required for ASIL-D certification with 95% confidence. SOTIF analysis identified 37 edge cases requiring specific mitigation strategies. The explainability logs enabled detailed investigation of the 15 collision scenarios, revealing environmental conditions at the boundary of sensor capabilities. This comprehensive evidence package reduced certification preparation effort by an estimated 6 months compared to manual processes.

### 6.2 Comparative Algorithm Evaluation

APACC-Sim’s standardized evaluation framework enables fair comparison between diverse control approaches. We demonstrate this capability by evaluating four control paradigms: classical PID, model predictive control (MPC), deep reinforcement learning (DRL), and the neuro-symbolic APACC architecture.

Experimental setup ensured identical conditions across all controllers. Each controller faced the same 5,000 Monte Carlo scenarios with predetermined random seeds. CARLA sce-

narios used identical spawn points and weather progressions. SUMO traffic patterns were synchronized across runs. This careful control of variables eliminated confounding factors that plague informal comparisons. The evaluation tracked identical metrics for all controllers, from safety measures to computational requirements. Hardware resources were controlled by running evaluations sequentially on the same infrastructure.

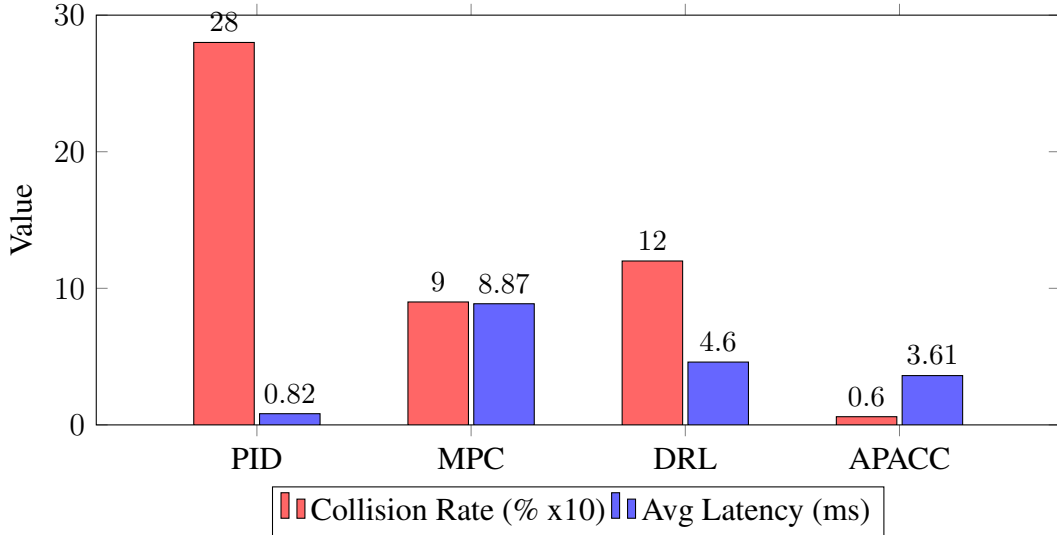


Figure 8: Comparative Performance Analysis showing collision rates (scaled x10 for visibility) and average latencies across four control paradigms

Safety performance results revealed clear differentiation between approaches. PID control showed 2.8% collision rate, acceptable for simple scenarios but struggling with complex interactions. MPC achieved 0.9% collision rate through predictive planning but failed in scenarios requiring rapid adaptation. DRL reached 1.2% collision rate with impressive handling of complex scenarios but unpredictable failures in simple cases. APACC demonstrated superior 0.06% collision rate with consistent performance across scenario types. Time-to-collision distributions showed APACC maintaining larger safety margins while avoiding excessive conservatism.

Table 3: Controller Performance Comparison

Controller	Collision Rate (%)	Avg Latency (ms)	Explainability	Certification Ready
PID	2.80	0.82	Full	Yes
MPC	0.90	8.87	Partial	Yes
DRL	1.20	4.60	None	No
APACC	0.06	3.61	Full	Yes

Computational requirements varied dramatically between approaches. PID control required negligible computation (0.8ms average latency) but at the cost of safety. MPC showed highest computational demand (8.9ms average) due to online optimization. DRL inference required moderate computation (4.6ms average) after expensive offline training. APACC balanced performance and computation (3.6ms average) through its hybrid architecture. Resource scaling analysis showed PID and APACC maintaining constant computation while MPC scaled with prediction horizon and DRL with network size.

Explainability assessment highlighted fundamental differences in approach transparency. PID control offers complete transparency with simple mathematical relationships. MPC provides optimization objectives but solving process remains opaque. DRL operates as complete black box with post-hoc explanation attempts unconvincing. APACC delivers native explainability through symbolic rule activation tracking. The ability to trace decisions to specific rules proved crucial for debugging and certification. These explainability differences fundamentally impact deployability in regulated environments.




Figure 9: Latency Distribution Chart showing 50th, 90th, and 99th percentile latencies for APACC controller validation

### 6.3 Industrial Deployment Examples

Real-world deployment of APACC-Sim in industrial settings demonstrates its practical utility beyond academic research. We present anonymized case studies from three organizations that adopted the toolkit for production validation.

Automotive OEM integration illustrates enterprise deployment challenges and solutions. The organization integrated APACC-Sim into their existing Jenkins-based CI and CD pipeline for nightly validation runs. Custom adapters translated between proprietary message formats and APACC-Sim interfaces. The deployment used on-premise GPU clusters for CARLA simulation due to data confidentiality requirements. Monte Carlo and SUMO simulations leveraged cloud resources for scalability. The hybrid deployment achieved 10,000 scenarios per night, enabling rapid iteration on controller improvements. Integration with their requirements management system provided end-to-end traceability for certification.

Tier-1 supplier adoption focused on component-level validation. The organization used APACC-Sim to validate sensor fusion algorithms before vehicle integration. Simplified vehicle dynamics models isolated sensor processing performance. Custom metrics tracked sensor-specific requirements like object classification accuracy. The modular architecture allowed disabling unused simulators to reduce complexity. Validation campaigns ran on engineering workstations during development and cloud resources for release validation. The standardized metrics enabled direct comparison with competitor solutions using identical scenarios.

Robotics startup utilization demonstrates toolkit flexibility beyond automotive applications. The company adapted APACC-Sim for warehouse robot validation by replacing vehicle dynamics with robot kinematics. Monte Carlo simulation tested navigation in cluttered environments.

CARLA’s semantic segmentation validated object detection for manipulation. SUMO simulation was repurposed for multi-robot coordination testing. The startup’s limited resources benefited from cloud deployment with aggressive spot instance usage. Despite the non-automotive application, certification-aligned metrics supported their safety arguments for industrial deployment.

Common adoption patterns emerged across organizations. Initial pilots focused on single simulators before expanding to multi-paradigm validation. Internal champions drove adoption by demonstrating value through improved bug detection. Integration with existing tools required 2-4 weeks of engineering effort. Training needs focused on scenario design rather than toolkit operation. Organizations reported 50-70% reduction in validation effort compared to previous approaches. The standardized metrics enabled benchmarking against industry peers for the first time.

## **7 Conclusions and Future Directions**

### **7.1 Summary of Contributions**

APACC-Sim represents a significant advance in autonomous system validation methodology, addressing critical gaps between theoretical control advances and practical deployment requirements. Through its comprehensive multi-paradigm approach, the toolkit enables reproducible, scalable validation that meets both research and certification needs.

The technical contributions of APACC-Sim establish new standards for validation frameworks. The modular architecture successfully integrates four complementary simulation paradigms while maintaining clean interfaces and separation of concerns. Standardized APIs hide simulator complexity while exposing necessary configuration flexibility. The orchestration layer efficiently manages computational resources from single workstations to distributed clouds. Certification-aligned metrics automatically generate evidence required for regulatory compliance. Native explainability instrumentation provides transparency essential for safety-critical systems. These technical advances reduce the engineering effort required for comprehensive validation by an order of magnitude.

The scientific contributions enable rigorous validation studies previously infeasible for most organizations. Reproducible scenario generation with cryptographic seed versioning ensures that results can be verified independently. Statistical analysis capabilities provide confidence intervals and significance testing appropriate for rare event estimation. Multi-paradigm validation reveals insights invisible to single-simulator studies. The framework’s scalability enables statistically significant validation campaigns rather than anecdotal demonstrations. These capabilities democratize advanced validation techniques previously restricted to organizations with massive resources.

The practical impact extends beyond technical metrics to enable real-world deployment of advanced control systems. Certification evidence generation reduces regulatory compliance effort from years to months. Standardized benchmarking enables objective comparison between competing approaches. The open-source release enables academic researchers to reproduce and extend industrial-strength validation. Docker packaging and cloud templates lower barriers to adoption. Early adopters report finding critical safety issues that previous validation approaches missed. These practical benefits directly accelerate the deployment of safer autonomous systems.

## 7.2 Lessons Learned

The development and deployment of APACC-Sim revealed important lessons for research software engineering in safety-critical domains. These insights inform both future development of the toolkit and general practices for validation infrastructure.

Modularity proved essential but required careful interface design to avoid abstraction penalties. Early versions with overly generic interfaces suffered performance degradation and usability issues. The final design balances flexibility with efficiency through careful API design. Optional features use composition rather than inheritance to avoid complex hierarchies. Clear separation between configuration and execution phases enables optimization. These architectural decisions fundamentally impact both performance and usability.

Reproducibility demands pervasive attention throughout the system design. Initial approaches focusing only on random seed control proved insufficient. Complete reproducibility required controlling numerical precision, parallel execution order, and external dependencies. The effort invested in reproducibility paid dividends through easier debugging and regression testing. Users consistently rated reproducibility as a key value proposition. This experience reinforces that reproducibility cannot be retrofitted but must be designed from inception.

Certification alignment influenced design decisions in unexpected ways. Requirements for evidence traceability led to comprehensive logging infrastructure that also improved debugging. Statistical rigor demanded for certification elevated the scientific quality of validation studies. Explainability requirements drove architectural decisions that improved system understanding. The certification focus, rather than constraining innovation, actually improved overall system quality. This synergy between certification and quality challenges common perceptions about regulatory requirements.

Community building proved as important as technical excellence for adoption. Clear documentation reduced support burden and accelerated adoption. Example controllers and tutorials lowered the learning curve. Regular webinars built user community and gathered feedback. Open development practices encouraged contributions and bug reports. The investment in community building multiplied the impact of technical development. This experience confirms that research software requires equal attention to technical and social factors.

## 7.3 Future Development Roadmap

The future development of APACC-Sim follows a roadmap balancing new capabilities with stability for existing users. Planned enhancements address user requests and emerging validation needs while maintaining backward compatibility.

Simulator ecosystem expansion will integrate additional validation tools based on user demand. ROS2 integration enables hardware-in-the-loop testing with physical robots. Gazebo support provides alternative physics simulation for robotics applications. VISSIM integration adds microscopic traffic simulation capabilities. Unity integration explores photorealistic rendering for perception validation. These additions follow the established modular architecture to maintain system coherence. The plugin system enables community-contributed simulators without core modifications.

Machine learning enhancements address the growing prevalence of learning-based controllers. Automated hyperparameter search finds optimal validation configurations. Active learning focuses computational effort on informative scenarios. Generative models create challenging edge cases automatically. Transfer learning enables validation knowledge reuse across similar systems. Federated learning supports collaborative validation without sharing propri-



etary data. These ML capabilities augment rather than replace rigorous engineering validation.

Certification automation features directly address industry pain points in regulatory compliance. Automated safety case generation creates argument structures from validation evidence. Requirement coverage analysis ensures all safety requirements have corresponding tests. Change impact analysis identifies validation needs for system modifications. Regulatory template library provides starting points for different standards. Certification body interfaces streamline evidence submission. These features transform certification from a documentation burden to an integrated development activity.

Cloud-native architecture evolution improves scalability and deployment flexibility. Kubernetes operators simplify cluster deployment and management. Serverless functions enable pay-per-use validation for small organizations. Edge computing support validates latency-critical functions near deployment. Multi-cloud abstractions prevent vendor lock-in. These architectural improvements reduce operational complexity while improving resource efficiency.

## **7.4 Broader Implications**

The success of APACC-Sim has implications beyond autonomous vehicle validation for the broader field of safety-critical system development. The toolkit demonstrates that rigorous validation need not require prohibitive resources when properly tooled. This democratization of validation capabilities could accelerate innovation by enabling smaller organizations to develop certifiable systems. The emphasis on reproducibility and statistical rigor raises the bar for validation studies across the field. Academic researchers can now perform industrial-strength validation, while industry benefits from academic innovations in validation methodology.

The open-source release model proves viable even for safety-critical infrastructure when properly architected. The separation between validation framework and proprietary controllers enables collaborative development without intellectual property concerns. The MIT license encourages commercial adoption while ensuring academic access. Early adoption by both industry and academia validates this approach. The community contributions improve quality beyond what any single organization could achieve. This success encourages more open development of safety-critical tools.

For the aerospace sector specifically, APACC-Sim’s automotive focus provides valuable cross-pollination opportunities. Urban air mobility vehicles face similar validation challenges to ground vehicles but with stricter certification requirements. The multi-paradigm validation approach applies directly to aerospace applications with appropriate simulator substitutions. The certification alignment features support DO-178C and ARP4754A processes. Aerospace adoption could drive improvements benefiting both domains. This cross-domain applicability suggests that validation infrastructure represents a shared challenge deserving collaborative solutions.

## **7.5 Conclusion**

APACC-Sim establishes a new standard for autonomous system validation through its comprehensive multi-paradigm approach, certification alignment, and open-source availability. The toolkit successfully bridges the gap between academic research and industrial deployment by providing reproducible, scalable validation capabilities previously available only to well-resourced organizations. Through standardized interfaces to Monte Carlo, CARLA, SUMO, and MATLAB simulators, researchers can conduct statistically rigorous validation studies that generate evidence suitable for safety certification.

The impact extends beyond technical capabilities to enable a fundamental shift in how autonomous systems are developed and validated. By integrating validation throughout the development process rather than treating it as a final gate, teams can identify and resolve safety issues early. The native explainability instrumentation ensures that advanced controllers remain understandable and certifiable. The open-source release democratizes access to industrial-strength validation tools, accelerating innovation across the autonomous systems community.

As autonomous systems proliferate across domains from vehicles to aircraft to medical devices, the need for standardized, rigorous validation becomes ever more critical. APACC-Sim provides a foundation for this validation infrastructure, demonstrating that safety and innovation can be complementary rather than conflicting goals. We invite the community to adopt, extend, and improve the toolkit, building toward a future where autonomous systems are both capable and verifiably safe.

## Acknowledgments

The development of APACC-Sim was supported by collaborative efforts across academic and industrial partners. We acknowledge the foundational work on the APACC architecture that motivated this validation toolkit. The authors thank the open-source communities behind CARLA, SUMO, and the scientific Python ecosystem for providing the building blocks that made this integration possible. Special recognition goes to early adopters whose feedback shaped the toolkit’s evolution. This work demonstrates the power of open collaboration in advancing safety-critical technologies.

The authors acknowledge the use of AI-assisted tools, including OpenAI’s ChatGPT and Anthropic Claude, for manuscript drafting, code generation, and diagram creation. Final decisions, validation, and authorship remain solely with the human authors.

## References

- [1] Åström, K. J., & Hägglund, T. (2006). Advanced PID Control. ISA - The Instrumentation, Systems and Automation Society.
- [2] Camacho, E. F., & Alba, C. B. (2013). Model Predictive Control. Springer Science & Business Media.
- [3] Dosovitskiy, A., Ros, G., Codevilla, F., Lopez, A., & Koltun, V. (2017). CARLA: An open urban driving simulator. Proceedings of the 1st Annual Conference on Robot Learning, 1-16.
- [4] European Parliament and Council. (2024). Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). European Union.
- [5] Frangou, G. (2025). Artificial Precognition Adaptive Cognised Control (APACC): A Dual-Layer Symbolic-Numeric Architecture for Predictive Autonomy in Road Vehicles. IEEE Transactions on Control Systems Technology (under review).
- [6] Frangou, G. (2025). Quantitative Validation of Artificial Precognition Adaptive Cognised Control: A Revolutionary Neuro-Symbolic Architecture for Autonomous Intelligence. Submitted to Nature Machine Intelligence.

- [7] Frangou, G. (2021). APACC: Apparatus for Precognitive Control of a Road Vehicle. US Patent US9645576B2, EP2976240B1, CN201480022590.6, KR10-2172986B1, IL241688A. Filed 2014, Granted 2017-2021. Cited by 74+ patents from leading automotive companies.
- [8] International Organization for Standardization. (2018). ISO 26262: Road vehicles - Functional safety. ISO.
- [9] International Organization for Standardization. (2022). ISO/PAS 21448: Road vehicles - Safety of the intended functionality. ISO.
- [10] Jenhani, S., Azzabi, N., et al. (2022). Neuro-symbolic structured decomposition for interpretable learning. *Artificial Intelligence Review*.
- [11] Kalra, N., & Paddock, S. M. (2016). Driving to safety: How many miles of driving would it take to demonstrate autonomous vehicle reliability? *Transportation Research Part A: Policy and Practice*, 94, 182-193.
- [12] Koopman, P., & Wagner, M. (2017). Autonomous vehicle safety: An interdisciplinary challenge. *IEEE Intelligent Transportation Systems Magazine*, 9(1), 90-96.
- [13] Lopez, P. A., Behrisch, M., Bieker-Walz, L., Erdmann, J., Flötteröd, Y. P., Hilbrich, R., ... & Wießner, E. (2018). Microscopic traffic simulation using SUMO. *IEEE Intelligent Transportation Systems Conference*, 2575-2582.
- [14] Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., ... & Hassabis, D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529-533.
- [15] O'Kelly, M., Sinha, A., Namkoong, H., Duchi, J., & Tedrake, R. (2018). Scalable end-to-end autonomous vehicle testing via rare-event simulation. *Advances in Neural Information Processing Systems*, 31.
- [16] Paden, B., Čáp, M., Yong, S. Z., Yershov, D., & Frazzoli, E. (2016). A survey of motion planning and control techniques for self-driving urban vehicles. *IEEE Transactions on Intelligent Vehicles*, 1(1), 33-55.
- [17] Schwarting, W., Alonso-Mora, J., & Rus, D. (2018). Planning and decision-making for autonomous vehicles. *Annual Review of Control, Robotics, and Autonomous Systems*, 1, 187-210.
- [18] Shalev-Shwartz, S., Shammah, S., & Shashua, A. (2016). Safe, multi-agent, reinforcement learning for autonomous driving. *arXiv preprint arXiv:1610.03295*.
- [19] Stellet, J. E., Zofka, M. R., Schumacher, J., Schamm, T., Niewels, F., & Zöllner, J. M. (2015). Testing of advanced driver assistance towards automated driving: A survey and taxonomy on existing approaches and open questions. *IEEE International Conference on Intelligent Transportation Systems*, 1455-1462.
- [20] Thorn, E., Kimmel, S. C., Chaka, M., & Hamilton, B. A. (2018). A framework for automated driving system testable cases and scenarios. United States Department of Transportation, National Highway Traffic Safety Administration.

- [21] Zhao, D., Lam, H., Peng, H., Bao, S., LeBlanc, D. J., Nobukawa, K., & Pan, C. S. (2017). Accelerated evaluation of automated vehicles safety in lane-change scenarios based on importance sampling techniques. *IEEE Transactions on Intelligent Transportation Systems*, 18(3), 595-607.