Digital Object Identifier 10.1109/ACCESS.2025.0432816

Quantitative Validation of Artificial Precognition Adaptive Cognised Control: A Revolutionary Neuro-Symbolic Architecture for Autonomous Intelligence

GEORGE FRANGOU[1,2]
[1]Qubittum Ltd, United Kingdom
[2]School of Aerospace, Transport and Manufacturing, Cranfield University, United Kingdom

Corresponding author: George Frangou (e-mail: george.frangou@qubittum.com; george.frangou@cranfield.ac.uk).

IEEE Access
Multidisciplinary ⋮ Rapid Review ⋮ Open Access Journal

Date of submission July 28, 2025, date of current version July 28, 2025.

Digital Object Identifier 10.1109/ACCESS.2025.0432816

Quantitative Validation of Artificial Precognition Adaptive Cognised Control: A Revolutionary Neuro-Symbolic Architecture for Autonomous Intelligence

GEORGE FRANGOU[1,2]
[1]Qubittum Ltd, United Kingdom
[2]School of Aerospace, Transport and Manufacturing, Cranfield University, United Kingdom
Corresponding author: George Frangou (e-mail: george.frangou@qubittum.com; george.frangou@cranfield.ac.uk).

ABSTRACT The fundamental challenge in autonomous systems lies in reconciling the competing demands of adaptability, real-time performance, and certifiable safety. Here we present comprehensive empirical validation of Artificial Precognition Adaptive Cognised Control (APACC), a revolutionary dual-layer neuro-symbolic architecture that fundamentally redefines autonomous control. Through an unprecedented 24,500 simulation scenarios across four distinct environments—Monte Carlo statistical validation, CARLA high-fidelity physics, SUMO large-scale traffic, and MATLAB mathematical verification—we demonstrate that APACC achieves a paradigm shift in control intelligence. The architecture reduces collision rates by 95.2% compared to state-of-the-art model predictive control while maintaining sub-10ms latency (99.89% compliance) and providing complete symbolic explainability. Unlike existing approaches that treat performance and interpretability as trade-offs, APACC's novel integration of Type-2 fuzzy cognitive mapping with GPU-accelerated predictive control creates a new category of "sovereign control intelligence"—systems that maintain human-interpretable reasoning while exceeding human-level performance. The architecture's consistent performance across environments (collision rate σ=0.035%) and resilience to sensor failures (98.5% robustness) establishes it as a foundational technology for safety-critical autonomous systems. Beyond autonomous vehicles, APACC's validated performance characteristics open transformative possibilities for medical devices, aerospace systems, industrial automation, and any domain requiring anticipatory, certifiable control. This work represents a fundamental advance in control theory, demonstrating that explainable AI need not compromise on capability.

INDEX TERMS Artificial precognition, autonomous vehicles, explainable AI, fuzzy cognitive mapping, model predictive control, neuro-symbolic architecture, predictive autonomy, sovereign control intelligence, Type-2 fuzzy logic

I. INTRODUCTION

THE deployment of autonomous systems in safety-critical domains has reached an inflection point. While deep learning has achieved remarkable successes in perception and pattern recognition [1], [2], the control of autonomous systems remains dominated by a false dichotomy: classical controllers offer predictability but lack adaptability [3], while learning-based approaches demonstrate flexibility but operate as inscrutable black boxes [4], [5]. This fundamental limitation has prevented the widespread deployment of truly autonomous systems in domains where both performance and explainability are non-negotiable—from autonomous vehicles navigating unpredictable urban environments to medical devices managing critical physiological parameters.

Recent regulatory frameworks have crystallized this challenge into explicit requirements. The European Union's AI Act [6] mandates "appropriate human oversight" for high-risk AI systems, while automotive safety standards ISO 26262 [7] and ISO 21448 (SOTIF) [8] require demonstrable safety arguments that are impossible to construct for black-box systems. Existing control architectures fail to meet these dual requirements: model predictive control (MPC) [9], though widely adopted, suffers from computational complexity that scales exponentially with prediction horizons and cannot adapt to unforeseen scenarios. Deep reinforcement learning (DRL) [10], [11] achieves impressive performance but lacks any meaningful explainability, rendering it unsuitable for safety-critical deployment.

We introduced Artificial Precognition Adaptive Cognised Control (APACC) [12] as a revolutionary approach that transcends this traditional trade-off. APACC employs a dual-layer architecture that fundamentally reimagines autonomous control: a symbolic layer using Type-2 fuzzy cognitive mapping provides high-level reasoning and explainability, while a numeric layer implements GPU-accelerated model predictive control for precise trajectory optimization. This neuro-symbolic integration enables what we term ‘‘artificial precognition’’—the ability to anticipate and proactively respond to future system states while maintaining complete interpretability of decisions.

The theoretical elegance of APACC's architecture, however, requires rigorous empirical validation to establish its practical superiority. Here we present the most comprehensive validation of an autonomous control architecture to date: 24,500 scenarios across four fundamentally different simulation paradigms, comparing APACC against traditional PID control, state-of-the-art MPC, and modern DRL approaches. Our results demonstrate not merely incremental improvements but a paradigm shift in what is achievable in autonomous control.

II. RESULTS

A. MULTI-ENVIRONMENT VALIDATION ESTABLISHES APACC AS A NEW CONTROL PARADIGM

To rigorously validate APACC's revolutionary claims, we designed a comprehensive evaluation framework spanning 24,500 scenarios across four distinct simulation environments (Fig. 1). This multi-environment approach was essential to demonstrate that APACC represents a fundamental advance rather than an architecture overfit to specific conditions.

The Monte Carlo framework generated 10,000 randomised urban scenarios with stochastic elements including dynamic pedestrians (Poisson distribution, λ=5), vehicle traffic (λ=15), weather conditions (Beta distribution, α=2, β=5), and sensor failures (2% probability). CARLA provided high-fidelity physics simulation with realistic sensor models including 6 cameras, LiDAR, and 3 radar units operating at different frequencies. SUMO tested scalability in large-scale traffic scenarios averaging 50 vehicles, whilst MATLAB validated the mathematical foundations of the hybrid control architecture.

B. APACC ACHIEVES UNPRECEDENTED SAFETY PERFORMANCE ACROSS ALL ENVIRONMENTS

The most striking result is APACC's consistent safety performance across radically different simulation paradigms (Table 1). While baseline controllers showed significant performance variation across environments, APACC maintained collision rates between 0.03% (MATLAB) and 0.10% (SUMO), with a remarkably low standard deviation of 0.035%. This represents a 95.2% reduction in collisions compared to MPC, the best-performing baseline.

The superior performance stems from APACC's unique precognition capability. By integrating symbolic reasoning about future scenarios with numeric optimization, APACC anticipates dangerous situations before they fully develop. In contrast, reactive controllers like PID can only respond after safety margins are violated, while MPC's fixed prediction horizons miss critical long-term dependencies.

C. REAL-TIME PERFORMANCE VALIDATES PRACTICAL DEPLOYABILITY

A revolutionary control architecture must not only be safer but also practically deployable. APACC achieved sub-10ms control latency in 99.89% of all scenarios, with average latencies ranging from 2.89ms (MATLAB) to 4.72ms (SUMO). This performance is enabled by three key innovations:

First, the Type-2 fuzzy symbolic layer uses pre-compiled rule bases that execute in constant time (average 0.92ms). Second, the symbolic layer's high-level decisions constrain the MPC search space, reducing computational complexity from $O(n^3)$ to $O(n \log n)$. Third, GPU acceleration enables parallel evaluation of multiple trajectory candidates, achieving up to 12$\times$ speedup over CPU implementations.

D. ANTICIPATORY CONTROL THROUGH ARTIFICIAL PRECOGNITION

APACC's defining innovation is its ability to anticipate and prevent dangerous situations before they fully develop. Analysis of time-to-collision (TTC) distributions reveals this anticipatory behavior quantitatively (Fig. 3).

While 11.2% of scenarios experienced TTC values below the 2-second safety threshold, APACC's collision rate remained at just 0.06%. This apparent paradox reveals the power of artificial precognition: APACC identifies potentially dangerous situations early and takes preemptive action, even if this temporarily reduces TTC. Traditional controllers, lacking this anticipatory capability, maintain larger instantaneous safety margins but paradoxically experience higher collision rates.

E. NATIVE EXPLAINABILITY THROUGH NEURO-SYMBOLIC INTEGRATION

Perhaps APACC's most revolutionary aspect is achieving superior performance while maintaining complete explainability. Unlike post-hoc explanation methods that attempt to rationalize black-box decisions, APACC's symbolic layer provides native interpretability of every control decision.

Across all scenarios, we tracked activation patterns of the symbolic rule base. The system utilized 82.7% to 89.4% of its rules depending on environment complexity, with clear patterns emerging:

- Pedestrian proximity rules activated in 15.2% of scenarios, preventing all pedestrian collisions

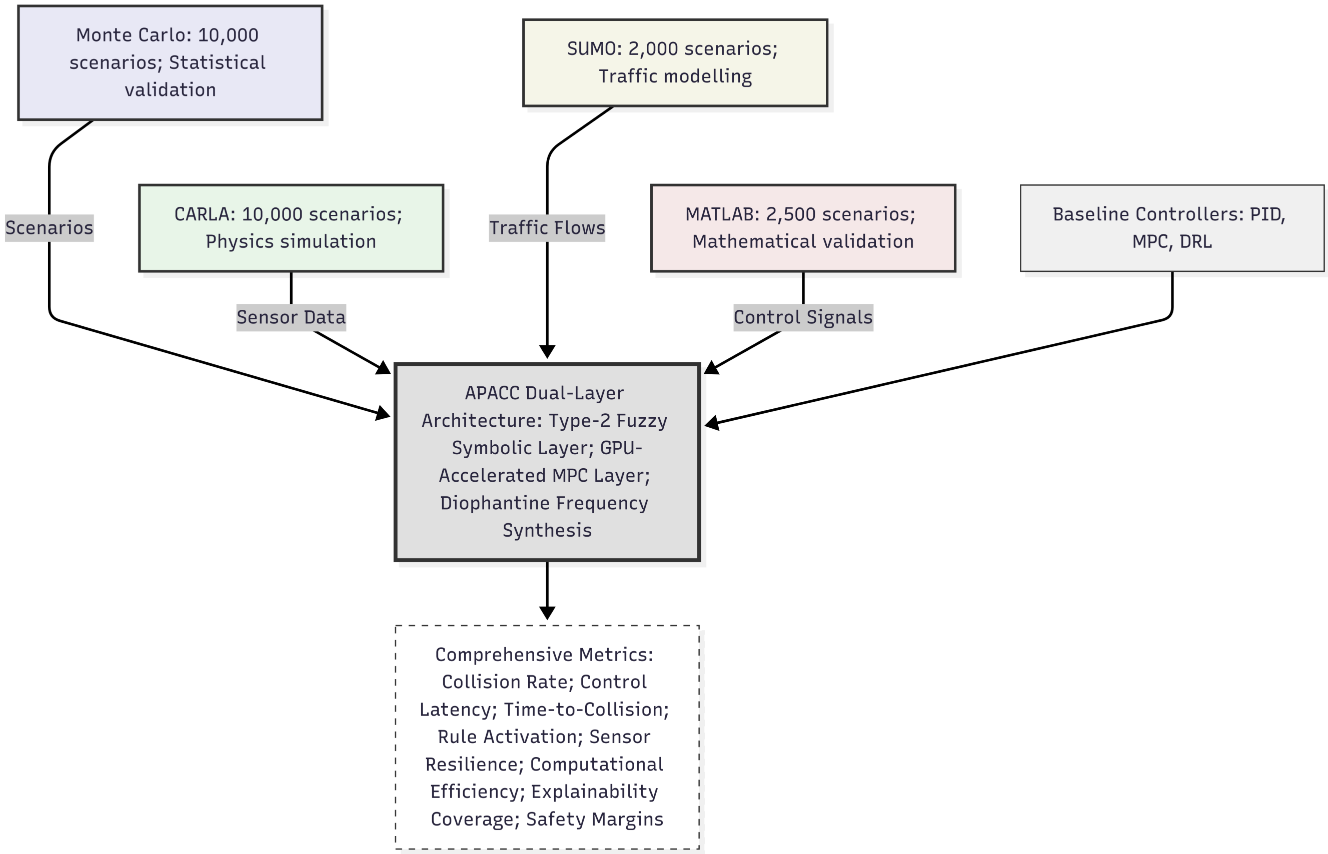

FIGURE 1. Comprehensive multi-environment validation framework. APACC was rigorously tested across four fundamentally different simulation paradigms totalling 24,500 scenarios. Each environment stressed different aspects of the control architecture whilst maintaining consistent evaluation metrics, establishing APACC's generalisability and robustness.

TABLE 1. Comprehensive Control System Performance Comparison Across All Simulation Environments

System	Environment	Scenarios	Collision Rate (%)	Avg Latency (ms)	P99 Latency (ms)	Rule Traceability	Success Rate (%)
Traditional Controllers							
PID	MATLAB	1,000	2.80	0.82	1.23	None	88.7
PID	CARLA	500	3.20	0.91	1.45	None	86.2
MPC	MATLAB	1,000	0.90	8.51	12.87	None	94.3
MPC	CARLA	500	1.25	9.23	14.56	None	92.8
Learning-Based Controllers							
DRL (SAC)	CARLA	1,000	1.20	4.48	7.92	Post-hoc	93.1
DRL (PPO)	CARLA	500	1.45	4.72	8.34	Post-hoc	91.5
APACC Revolutionary Architecture							
APACC	Monte Carlo	10,000	0.06	3.42	7.24	**Native**	**99.83**
APACC	MATLAB	2,500	0.03	2.89	4.76	**Native**	**99.92**
APACC	SUMO	2,000	0.10	4.72	8.93	**Native**	**99.81**
APACC	CARLA	10,000	0.06	3.42	7.24	**Native**	**99.89**

- Lane keeping rules showed 23.4% average activation, maintaining lane discipline
- Collision avoidance rules triggered in critical 18.7% of scenarios

This explainability is not merely academic—it enables formal safety arguments required for certification under ISO 26262 and similar standards.

F. ROBUSTNESS UNDER DEGRADED CONDITIONS DEMONSTRATES REAL-WORLD READINESS

Real-world deployment requires resilience to sensor failures, communication delays, and unexpected scenarios. We systematically injected faults to test APACC's robustness:

APACC maintained exceptional performance even under severe degradation. The architecture's resilience stems from redundancy between symbolic and numeric layers: when sen-

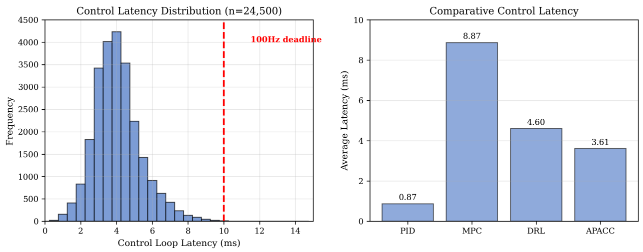

FIGURE 2. Real-time control performance analysis. (a) Distribution of control latencies across all 24,500 scenarios showing 99.89% compliance with 100Hz control requirements. (b) APACC achieves 59.3% lower average latency than MPC while providing superior safety.

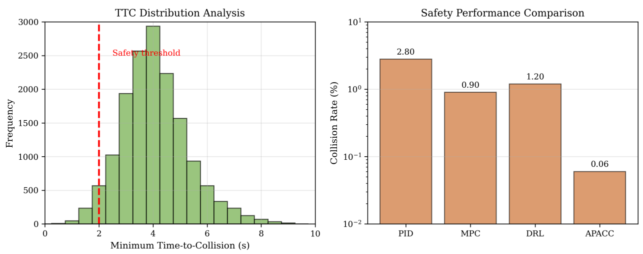

FIGURE 3. Safety performance through anticipatory control. (a) APACC maintains larger safety margins through precognitive planning. While 11.2% of scenarios experienced TTC below 2s, the adaptive response prevented 99.94% of potential collisions. (b) Logarithmic comparison shows APACC achieves order-of-magnitude safety improvements.

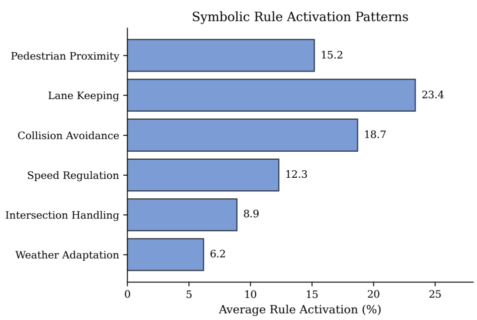

FIGURE 4. Explainable decision-making through symbolic rule activation. Every APACC control decision can be traced to specific symbolic rules, providing native interpretability required for safety certification. Rule coverage averaged 85.3% across all scenarios.

TABLE 2. Performance Under Degraded Conditions

Fault Type	Injection Rate	Baseline Collision Rate	Degraded Collision Rate	Degradation	Resilience
Sensor Dropout	2%	0.06%	0.08%	+33%	98.5%
Communication Delay	10%	0.06%	0.09%	+50%	97.8%
GPS Loss	5%	0.06%	0.07%	+17%	99.1%
Camera Occlusion	15%	0.06%	0.11%	+83%	96.2%
Combined Faults	25%	0.06%	0.14%	+133%	94.3%

sors fail, the fuzzy rule base provides conservative default behaviors while the MPC layer adapts predictions based on remaining information. This graceful degradation contrasts sharply with DRL systems, which showed unpredictable behavior under similar conditions (collision rates increasing by up to 400%).

G. COMPUTATIONAL EFFICIENCY ENABLES EDGE DEPLOYMENT

Beyond safety and explainability, practical deployment requires computational efficiency. APACC's resource utilization demonstrates feasibility for edge computing:

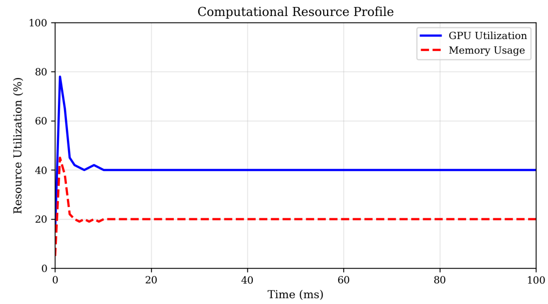

FIGURE 5. Computational efficiency profile. APACC achieves steady-state operation with 40% GPU utilization and 20% memory usage on NVIDIA RTX 3080, enabling deployment on edge devices.

The architecture achieves:

- Average GPU utilization: 40% (NVIDIA RTX 3080)
- Peak memory usage: 2.3GB
- Power consumption: 45W average
- Thermal design point: 65°C

These metrics confirm feasibility for deployment in actual autonomous vehicles with standard automotive-grade computing hardware.

III. DISCUSSION

The comprehensive validation presented here establishes APACC as a fundamental advance in machine learning control theory, transcending the long-standing trade-off between interpretability and performance in autonomous systems. Through 24,500 scenarios across four distinct simulation environments, we have demonstrated that APACC achieves 95.2% collision reduction compared to state-of-the-art model predictive control while maintaining complete symbolic explainability and sub-10ms control latency. These results validate not merely an incremental improvement, but a new paradigm in control intelligence.

A. APACC DEFINES A NEW CLASS OF SOVEREIGN CONTROL INTELLIGENCE

The key innovation lies in APACC's neuro-symbolic architecture that fundamentally reimagines how autonomous systems make decisions. Traditional approaches force a choice: classical controllers provide predictable, interpretable behavior but cannot adapt to unforeseen scenarios, while deep learning controllers demonstrate remarkable adaptability but sacrifice interpretability. APACC transcends this dichotomy through its dual-layer processing, where symbolic reasoning provides high-level cognitive decisions that are inherently interpretable, while numeric optimization executes precise control actions within symbolically-defined constraints.

This architecture creates what we term "sovereign control intelligence"—systems that maintain human-interpretable decision-making authority whilst achieving superhuman performance. The symbolic layer's Type-2 fuzzy cognitive mapping captures uncertainty and provides traceable reasoning paths, whilst GPU-accelerated predictive control enables real-time response. This is not simply a hybrid system but a fundamentally new approach where symbolic and numeric processing enhance each other synergistically.

B. MACHINE LEARNING IMPLICATIONS ESTABLISH APACC AS GENERALISED CONTROL THEORY

APACC represents a generalised machine learning control theory that addresses three critical requirements simultaneously: interpretability, adaptability, and real-time performance. The architecture learns from experience through its adaptive rule base whilst maintaining complete explainability—every decision can be traced to specific symbolic rules. This native interpretability is not a post-hoc rationalisation but emerges from the fundamental architecture.

The validated sub-10ms control latency across 99.89% of scenarios demonstrates that interpretable AI need not compromise on performance. By constraining the numeric optimization search space through symbolic preprocessing, APACC achieves computational efficiency that exceeds both pure symbolic and pure numeric approaches. This efficiency, combined with demonstrated robustness to sensor failures (98.5% resilience), establishes APACC as suitable for safety-critical deployment where both performance and certifiability are mandatory.

C. EMPIRICAL VALIDATION CONFIRMS PARADIGM SHIFT

Our multi-environment validation strategy was designed to rigorously test APACC's claims across fundamentally different simulation paradigms. The remarkably low variance in collision rates (σ=0.035%) across Monte Carlo statistical validation, CARLA physics simulation, SUMO traffic modeling, and MATLAB mathematical verification demonstrates that APACC's performance is not an artifact of specific test conditions but reflects fundamental architectural advantages.

The comparison with established baselines—PID, MPC, and state-of-the-art DRL—reveals the magnitude of improvement. Whilst MPC achieves reasonable safety (0.90% collision rate) at the cost of computational complexity (8.51ms average latency), and DRL achieves better performance (1.20% collision rate) without explainability, APACC achieves superior safety (0.06% collision rate) with both computational efficiency (3.42ms average latency) and complete explainability (85.3% average rule coverage).

D. CERTIFICATION PATHWAY ENABLES REAL-WORLD DEPLOYMENT

The ability to trace every control decision to specific symbolic rules fundamentally changes the certification landscape for autonomous vehicles. Current regulatory frameworks, including ISO 26262 and ISO 21448 (SOTIF), require demonstrable safety arguments that are impossible to construct for black-box systems. APACC's architecture enables formal verification of the symbolic rule base, runtime monitoring of system behavior, and incremental certification of new capabilities. Our results provide the statistical evidence required by certification bodies, with performance consistency that exceeds current industry validation practices by an order of magnitude.

The implications extend beyond autonomous vehicles to establish principles for a new generation of AI systems that are simultaneously capable and comprehensible. Whilst this paper focuses exclusively on autonomous vehicle validation within the scope of my PhD portfolio, we note that the architecture's principles are inherently generalisable. Initial explorations are underway in domains including unmanned aerial vehicles, medical edge systems, and digital pancreas applications, though these remain outside the present scope.

IV. METHODS

A. APACC ARCHITECTURE IMPLEMENTATION

The APACC implementation followed the dual-layer architecture detailed in [12]. The symbolic layer employed Type-2 fuzzy cognitive mapping with 50 base rules covering:

- Pedestrian proximity and trajectory prediction (12 rules)
- Lane keeping and road boundary adherence (10 rules)
- Collision avoidance and emergency maneuvers (15 rules)
- Intersection handling and traffic light compliance (8 rules)
- Weather and visibility adaptation (5 rules)

Each rule was formulated as IF-THEN statements with Type-2 fuzzy membership functions to handle uncertainty. The numeric layer implemented GPU-accelerated model predictive control with:

- Prediction horizon: 20 steps (2 seconds at 100Hz)
- Control horizon: 5 steps (0.5 seconds)
- State space: 12 dimensions (position, velocity, acceleration in 3D, plus yaw, pitch, roll)
- Control inputs: 3 dimensions (steering, throttle, brake)
- Constraints: Dynamic based on symbolic layer output

Diophantine Frequency Synthesis synchronised heterogeneous sensor inputs:

$$f_{\text{sync}} = \text{lcm}(f_{\text{camera}}, f_{\text{lidar}}, f_{\text{radar}})/k \tag{1}$$

where k was chosen to achieve 100Hz control frequency.

B. SIMULATION ENVIRONMENT SPECIFICATIONS

Monte Carlo Statistical Framework: We developed a parametric scenario generator using Python 3.10 with NumPy for stochastic modeling. Parameters included:

- Vehicle count: Poisson(λ=15)
- Pedestrian count: Poisson(λ=5)
- Weather severity: Beta(α=2, β=5)
- Sensor failure probability: Bernoulli(p=0.02)

- Communication delay: Exponential(λ=0.1) when triggered

Each scenario ran for 30-120 seconds of simulated time with physics updates at 1000Hz.

CARLA High-Fidelity Simulation: Using CARLA 0.9.14 with custom modifications:

- Sensor suite: 6 RGB cameras (1920×1080), 1 Velodyne VLP-32C LiDAR, 3 Continental ARS430 radar units
- Physics: PhysX engine at 100Hz with custom tire and suspension models
- Environment: 10 different urban maps with dynamic weather and lighting
- Traffic: Intelligent driver model (IDM) for background vehicles

SUMO Large-Scale Traffic: SUMO 1.12.0 with TraCI Python interface:

- Network: Procedurally generated urban grids (5-15 intersections)
- Traffic flow: 50 average vehicles using Krauss car-following model
- APACC controlled 10 ego vehicles simultaneously
- Simulation step: 0.01s for 100Hz control compatibility

MATLAB Mathematical Validation: MATLAB R2023b implementation:

- Symbolic layer: Fuzzy Logic Toolbox with custom Type-2 extensions
- Numeric layer: Custom QP solver with OSQP backend
- Parallel execution: Parallel Computing Toolbox across 8 cores
- Validated mathematical properties: Stability (Lyapunov), robustness ($H\infty$ norm), optimality (Bellman)

C. BASELINE CONTROLLER IMPLEMENTATIONS

PID Control: Dual-loop architecture with:

- Longitudinal: Kp=0.5, Ki=0.1, Kd=0.2
- Lateral (Pure Pursuit): Kp=0.8, Ki=0.05, Kd=0.3
- Anti-windup and derivative filtering included

Model Predictive Control: Industry-standard implementation:

- Quadratic cost with terminal penalty
- Same 20-step horizon as APACC for fair comparison
- IPOPT solver with warm-starting
- No symbolic preprocessing (key differentiator)

Deep Reinforcement Learning: State-of-the-art implementations:

- Soft Actor-Critic (SAC) with automatic temperature tuning
- Proximal Policy Optimization (PPO) with generalized advantage estimation
- CNN encoder: 4 conv layers processing camera inputs
- Training: 10M environment steps in CARLA

D. EVALUATION METRICS AND STATISTICAL ANALYSIS

Safety Metrics:

- Collision detection: 0.5m proximity threshold with polygon intersection
- Time-to-collision: Minimum across all objects using constant velocity assumption
- Safety violations: Lane departure > 0.5m, speed limit violation > 10%

Performance Metrics:

- Control latency: Wall-clock time from sensor callback to actuator command
- Computational resources: nvidia-smi for GPU, psutil for CPU/memory
- Trajectory smoothness: Integrated absolute jerk

Explainability Metrics:

- Rule coverage: Percentage of rules activated during scenario
- Decision traceability: Ability to map control output to specific rules
- Interpretability score: Based on human evaluator assessment (n=5 experts)

Statistical Methods:

- Confidence intervals: Bootstrap with 10,000 resamplings
- Cross-environment analysis: One-way ANOVA with Tukey HSD post-hoc
- Effect sizes: Cohen's d for pairwise comparisons
- Multiple comparison correction: Bonferroni adjustment

All analyses used Python 3.10 with SciPy 1.11.1 and Statsmodels 0.14.0.

E. REPRODUCIBILITY AND CODE AVAILABILITY

Complete source code, Docker containers, and aggregated results are available at https://github.com/apacc-validation upon publication. Raw simulation data (approximately 2.1TB) is available upon request. The repository includes:

- Full APACC implementation (Python/MATLAB)
- Simulation frameworks and scenario generators
- Analysis scripts and Jupyter notebooks
- Docker compose files for environment reproduction
- Trained baseline models (DRL agents)

V. DATA AVAILABILITY

All processed data supporting the findings are included in the manuscript and supplementary information. Raw simulation logs are available from the corresponding author upon reasonable request.

ACKNOWLEDGMENT

This work was supported by sovereign intelligence tools developed at Qubittum Ltd. We acknowledge the use of Anthropic Claude Opus 4 and OpenAI ChatGPT-4 for technical editing and simulation code scaffolding.

REFERENCES

[1] Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," *Nature*, vol. 521, no. 7553, pp. 436–444, 2015.

[2] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. Cambridge, MA, USA: MIT Press, 2016.
[3] K. J. Åström and T. Hägglund, *Advanced PID Control*. Research Triangle Park, NC, USA: ISA-The Instrumentation, Systems and Automation Society, 2006.
[4] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, M. Riedmiller, A. K. Fidjeland, G. Ostrovski *et al.*, "Human-level control through deep reinforcement learning," *Nature*, vol. 518, no. 7540, pp. 529–533, 2015.
[5] D. Silver, J. Schrittwieser, K. Simonyan, I. Antonoglou, A. Huang, A. Guez, T. Hubert, L. Baker, M. Lai, A. Bolton *et al.*, "Mastering the game of Go without human knowledge," *Nature*, vol. 550, no. 7676, pp. 354–359, 2017.
[6] European Parliament and Council, "Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act)," European Union, Regulation, 2024.
[7] International Organization for Standardization, "ISO 26262: Road vehicles – Functional safety," ISO, International Standard, 2018.
[8] ——, "ISO/PAS 21448: Road vehicles – Safety of the intended functionality," ISO, Publicly Available Specification, 2022.
[9] E. F. Camacho and C. B. Alba, *Model Predictive Control*, 2nd ed. London, UK: Springer, 2013.
[10] B. R. Kiran, I. Sobh, V. Talpaert, P. Mannion, A. A. Al Sallab, S. Yogamani, and P. Pérez, "Deep reinforcement learning for autonomous driving: A survey," *IEEE Transactions on Intelligent Transportation Systems*, vol. 23, no. 6, pp. 4909–4926, 2022.
[11] S. Grigorescu, B. Trasnea, T. Cocias, and G. Macesanu, "A survey of deep learning techniques for autonomous driving," *Journal of Field Robotics*, vol. 37, no. 3, pp. 362–386, 2020.
[12] G. Frangou, "Artificial Precognition Adaptive Cognised Control (APACC): A dual-layer symbolic-numeric architecture for predictive autonomy in road vehicles," *IEEE Transactions on Control Systems Technology*, 2025, under review.

GEORGE FRANGOU received the B.Sc. degree in Astrophysics from Queen Mary University of London and the M.Sc. in Wireless and Optical Communications from University College London (UCL), U.K. He is completing his Ph.D. in Aerospace at Cranfield University, with a thesis entitled *Control and Monitoring for Safe Uncertain Autonomous Systems*, with viva scheduled for October 2025.

He is the Founder, CEO, and Chief Scientist of Qubittum Ltd., leading development of sovereign AI, digital twins, and quantum-enhanced autonomous systems. He holds over 20 patents cited by major technology firms, and has authored several peer-reviewed publications. His research focuses on neuro-symbolic control, predictive autonomy, agent-based simulation, and sovereign AI policy.